

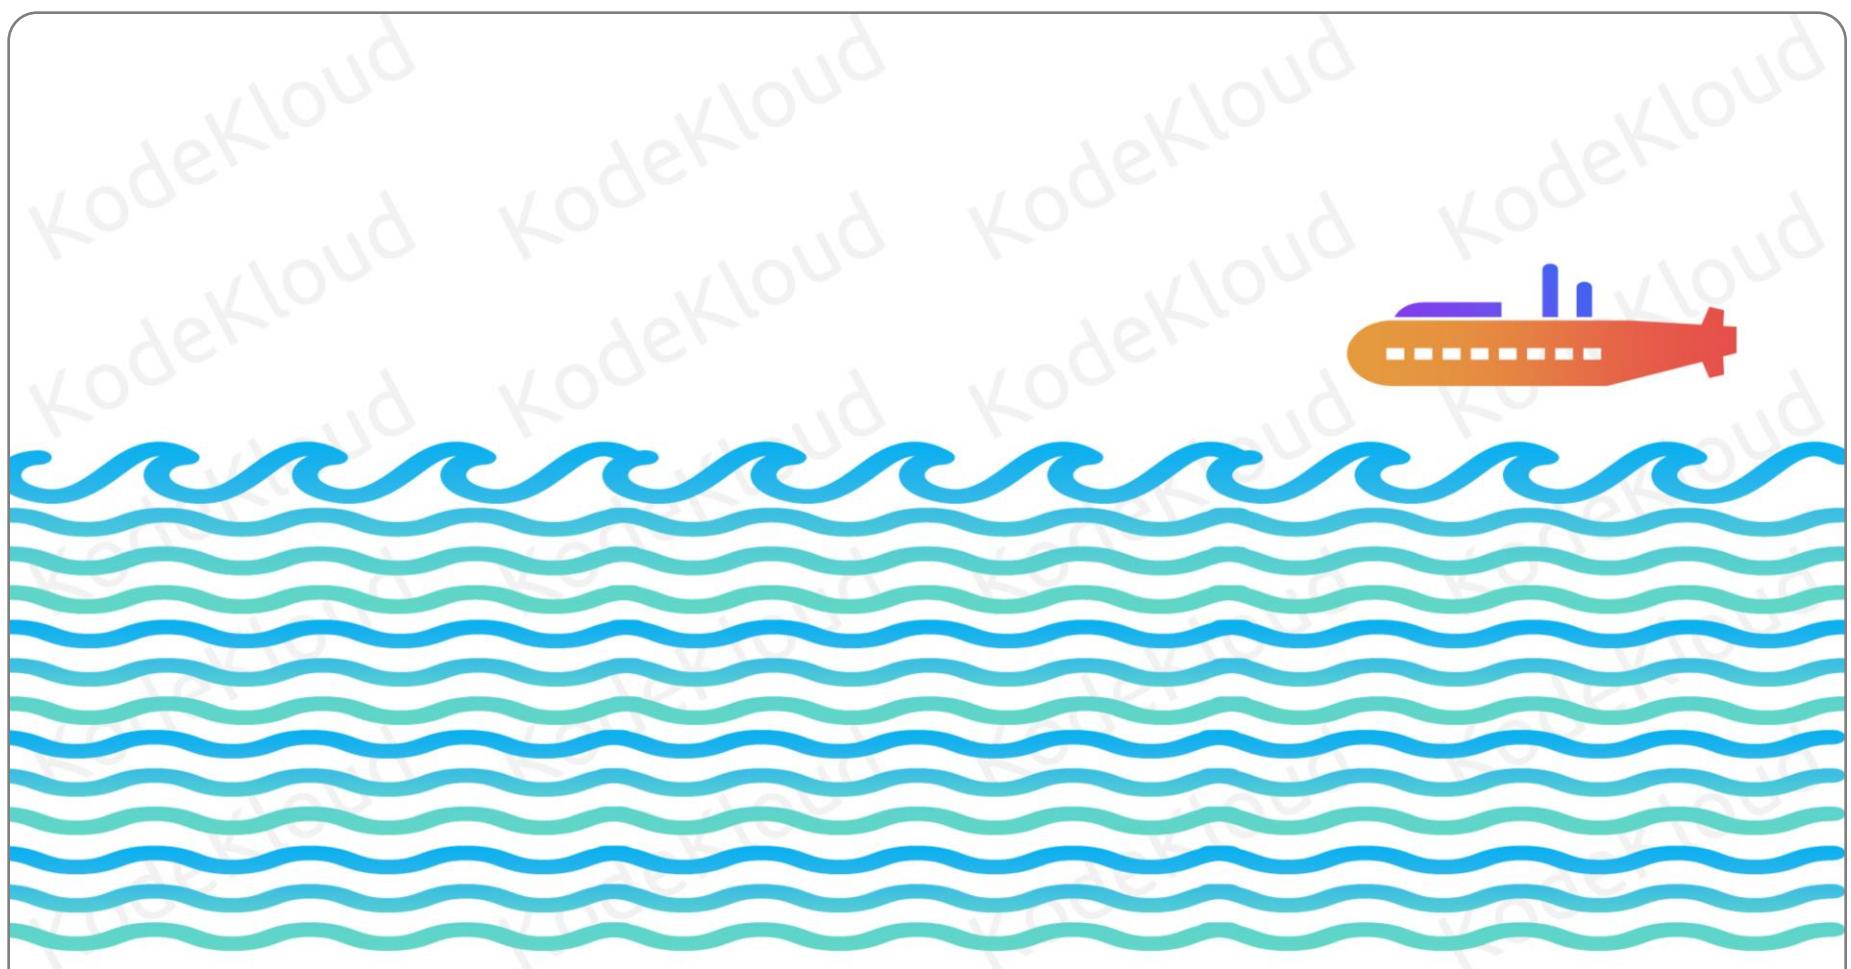


KodeKloud

© Copyright KodeKloud

Follow us on <https://kodekloud.com/> to learn more about us.

Elasticsearch FluentBit and Kibana – Introduction



© Copyright KodeKloud



Elasticsearch



Kibana



FluentBit



Mastering Elasticsearch Fundamentals



Elasticsearch

1 | Nodes

2 | Cluster

3 | Index

4 | Document

5 | CRUD Operations

6 | Mapping



Introduction to the EFK Stack



Kibana

- 1 | Kibana visualizations and dashboards
- 2 | Using Kibana's Discover interface for log exploration
- 3 | Kibana Query Language (KQL)





FluentBit

- 1 | Logstash's role within the ELK stack
- 2 | FluentBit vs FluentD
- 3 | Logstash vs FluentBit
- 4 | FluentBit Configurations





01

Setting up FluentBit to monitor K8 infra



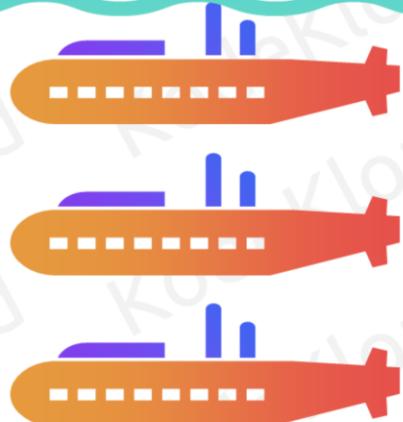
02

Monitoring a mock e-commerce app application with FluentBit and Elasticsearch



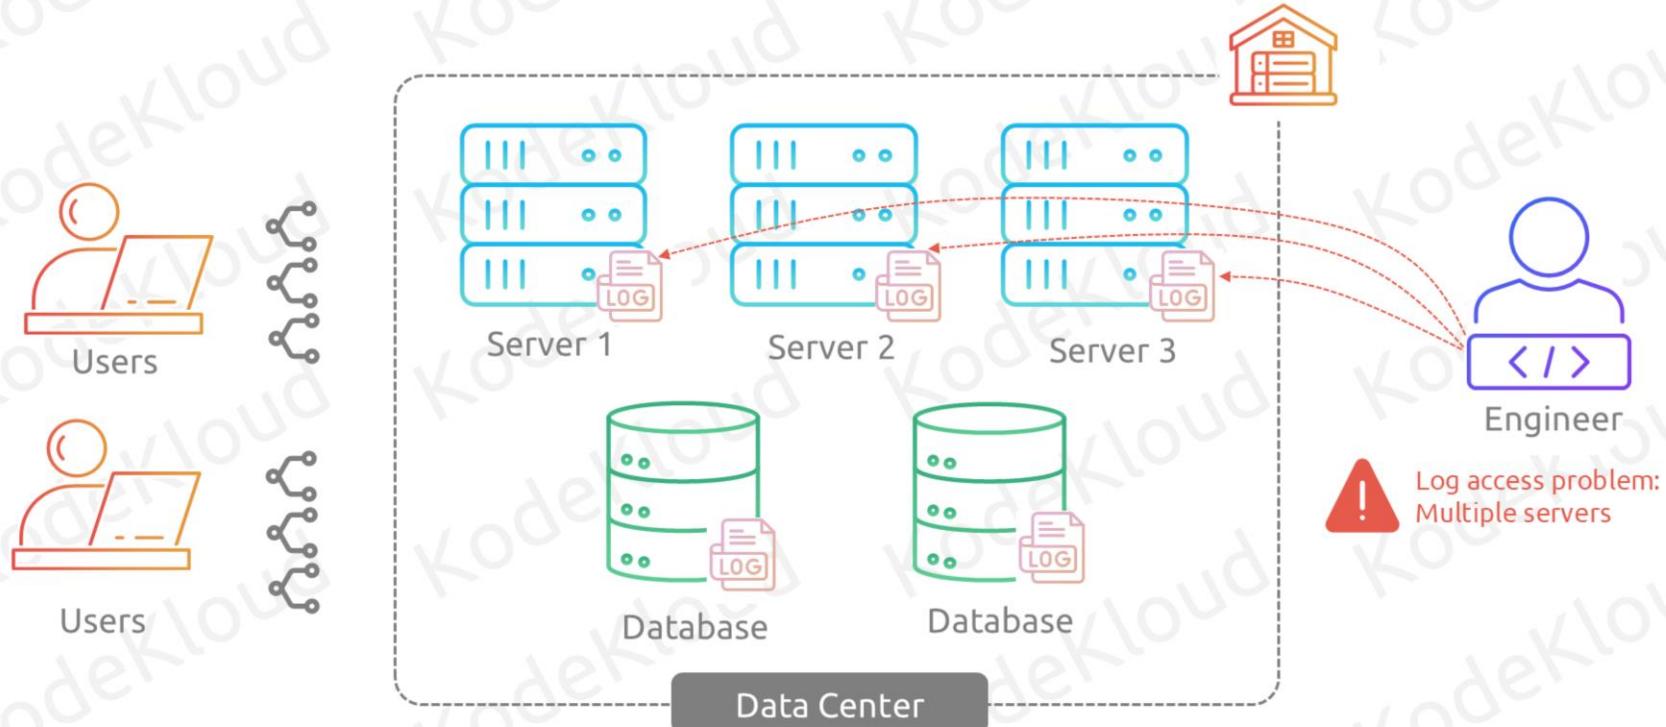
03

Monitoring a Python application with FluentBit and Elasticsearch

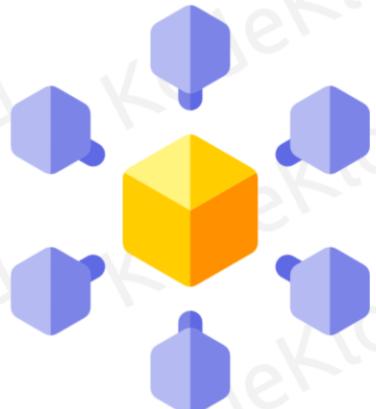


What Is Logging and Why Does It Matter?

Multi-Server Log Management Challenges



Microservices and Containerization

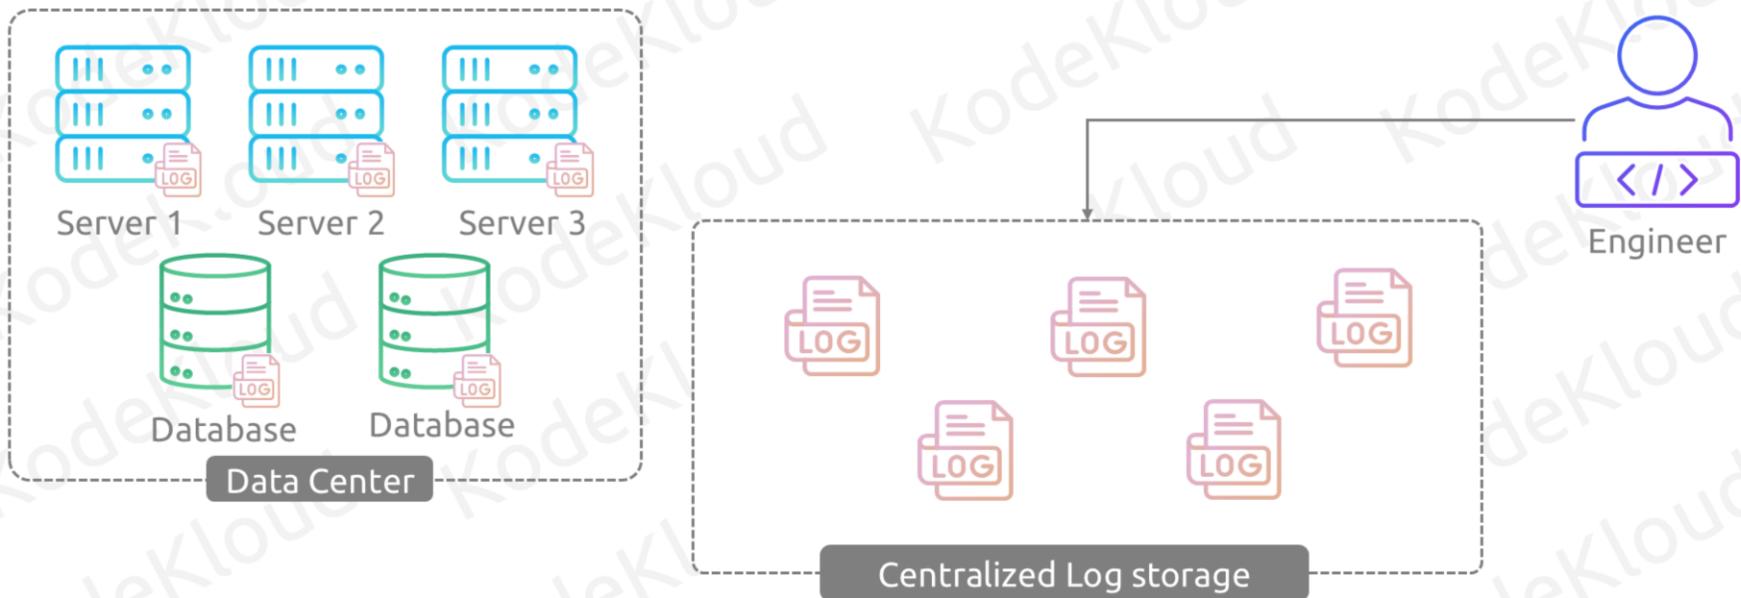


Microservices Architecture

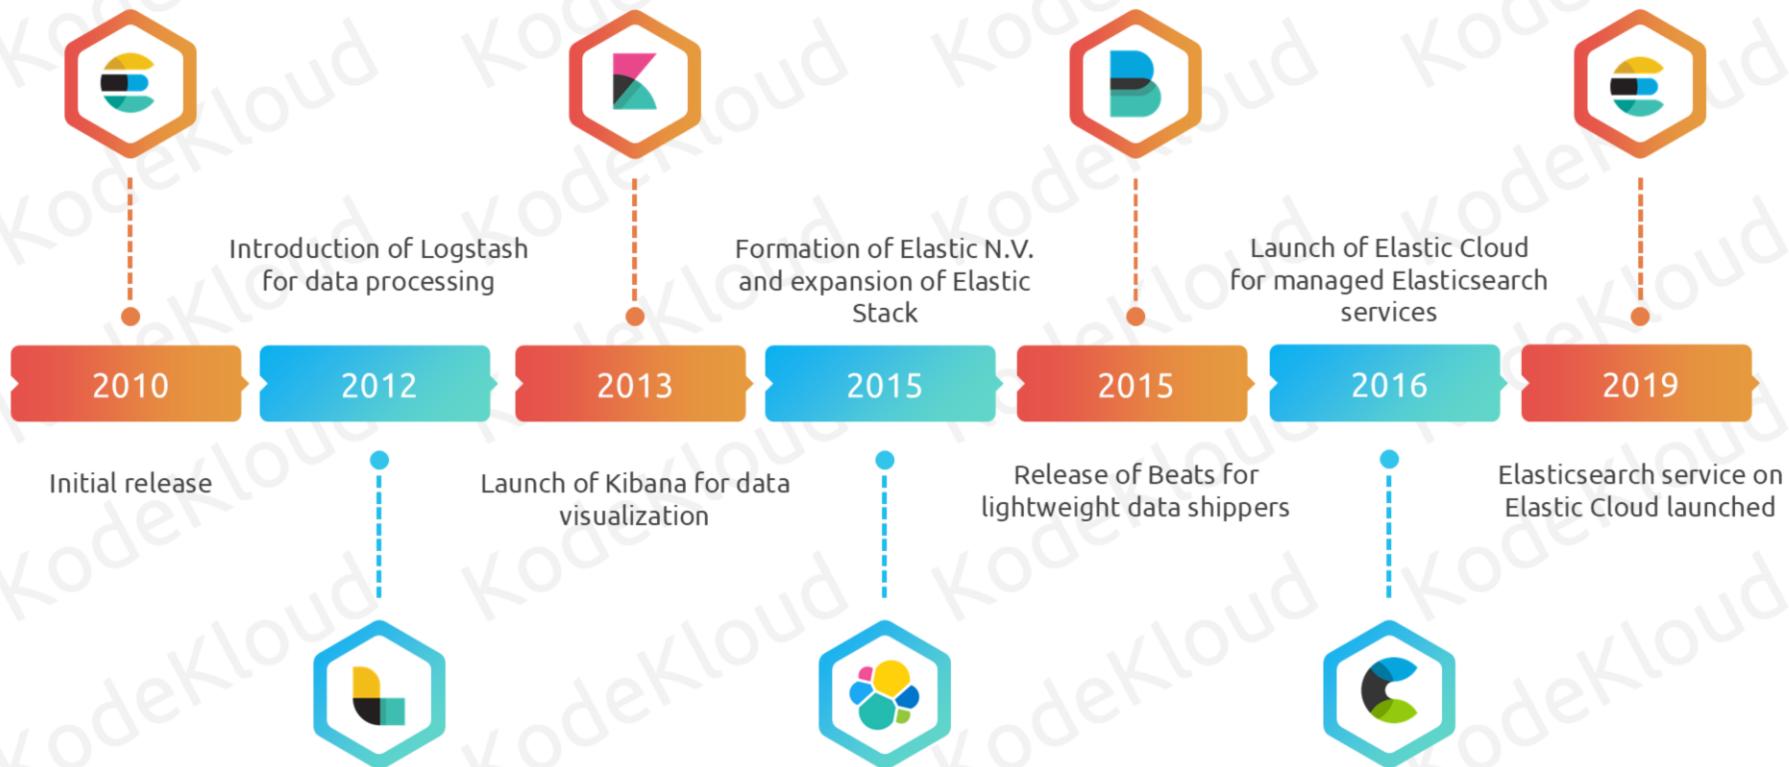


Docker

Imagining Centralized Logging

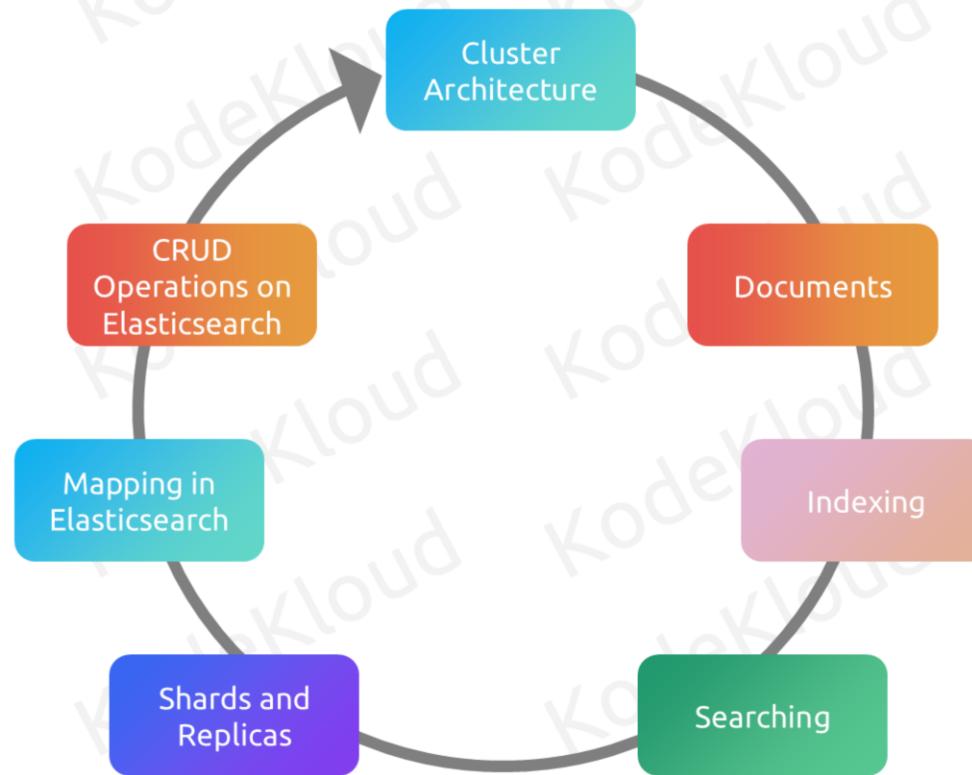


Elasticsearch and Its Evolution

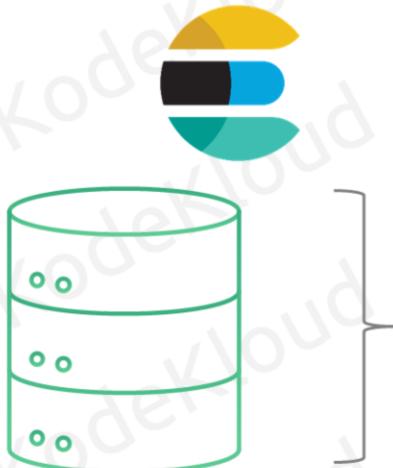


Mastering Fundamentals of Elasticsearch

Elasticsearch – What We Will Learn



Nodes and Cluster in Elasticsearch



This is a **single-node** elasticsearch setup

Recommended Node Configuration

Memory: 16GB

CPU: 8 Cores

Disk: SSDs

Network: At least 1 Gbps network connection

Nodes and Cluster in Elasticsearch



Nodes and Cluster in Elasticsearch



But do all nodes in the cluster play the same role ?

Nodes and Cluster in Elasticsearch

01



Collection of nodes

02



Single namespace

03



Redundancy and high availability

04



Scalability

NodeRoles in Elasticsearch



Master Node

Data Node

Data Content

Data Hot

Data Warm

Data Cold

Data Frozen

Data Ingest

ML Node

Transform

Remote Cluster
Client

NodeRoles in Elasticsearch

Master Node

Manages cluster changes and metadata

Data Node

Stores, indexes data; handles search and aggregation

Data Ingest

Pre-processes documents via ingest pipelines before indexing

ML Node

Handles machine learning tasks like anomaly detection

Transform

Executes data transformations, pivot, and latest transforms

© Copyright KodeKloud

Master Node-Manages the cluster by handling changes and maintaining cluster metadata.

Data Node-Stores and indexes data, handling search and aggregation requests

Data Ingest-Pre-processes documents before indexing using ingest pipelines.

ML Node-Performs machine learning tasks like anomaly detection within the cluster.

Transform-Executes data transformations and operations such as pivot and latest transforms.

NodeRoles in Elasticsearch

Remote Cluster Client

Acts as a gateway for cross-cluster searches

Data Cold

Stores infrequently accessed data, cost-optimized over performance

Data Frozen

Holds rarely accessed data, prioritizing cost efficiency with high latency

Data Hot

Stores frequently accessed data, optimized for low latency and high performance

Data Warm

Balances storage efficiency and performance for moderately accessed

© Copyright KodeKloud

Remote Cluster Client-Facilitates cross-cluster search by acting as a gateway to remote clusters.

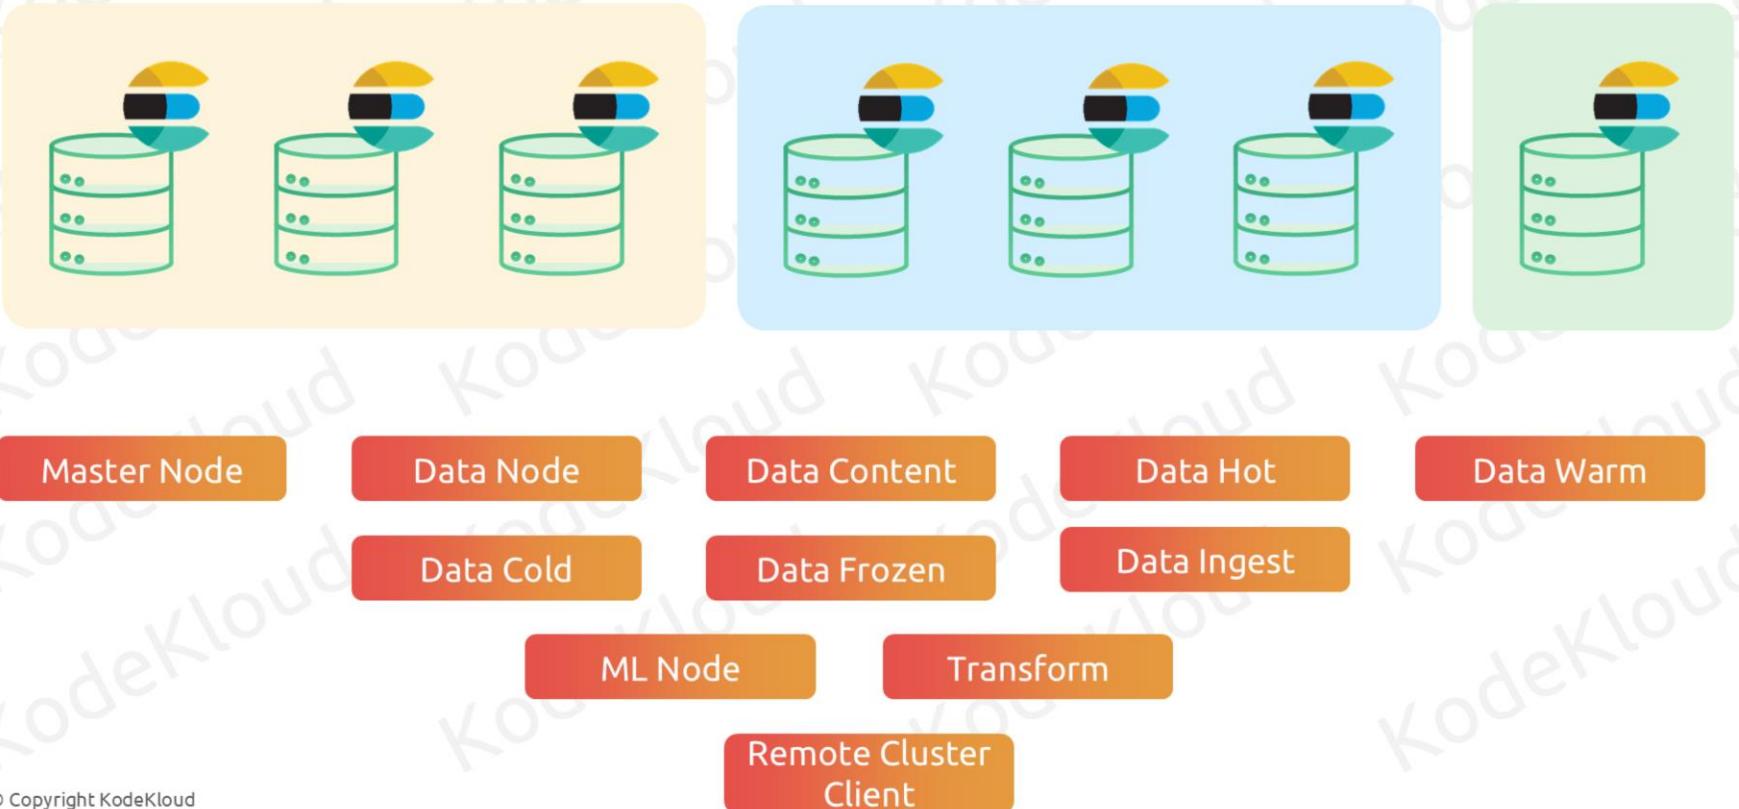
Data Cold-Contains infrequently accessed data, optimized for cost-effective storage over performance.

Data Frozen-Stores rarely accessed data, utilizing the most cost-efficient storage with the highest latency.

Data Hot-Stores frequently accessed and updated data, optimized for low latency and high performance.

Data Warm-Holds less frequently accessed data, balancing performance and storage efficiency.

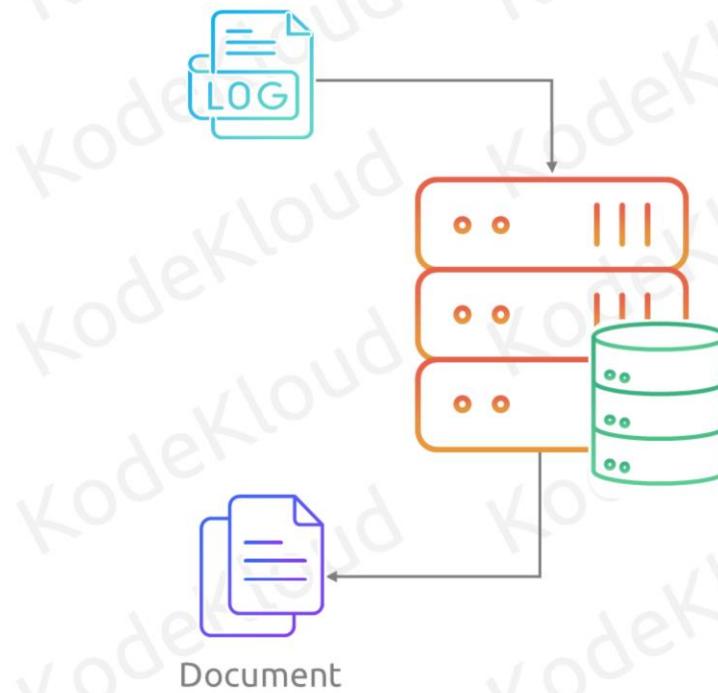
NodeRoles in Elasticsearch



So here I wanted to sort or move these roles to different servers and show how when having a big cluster of elasticsearch each node could be playing different roles.

Here I am using same server type but this can be different server icons

What Gets Stored in Elasticsearch?



© Copyright KodeKloud

In Elasticsearch, a document is a basic unit of information that can be indexed and stored within the platform. Documents are stored in a structured JSON format and are used to perform and optimize search operations in real-time

What Gets Stored in Elasticsearch?



Document

- 1 | Fundamental unit of data
- 2 | Stored in structured JSON
- 3 | Optimized for real-time search

Features of a Document in Elasticsearch

01

JSON Format



02

Indexing



03

Node



04

Sharding and Replicas



05

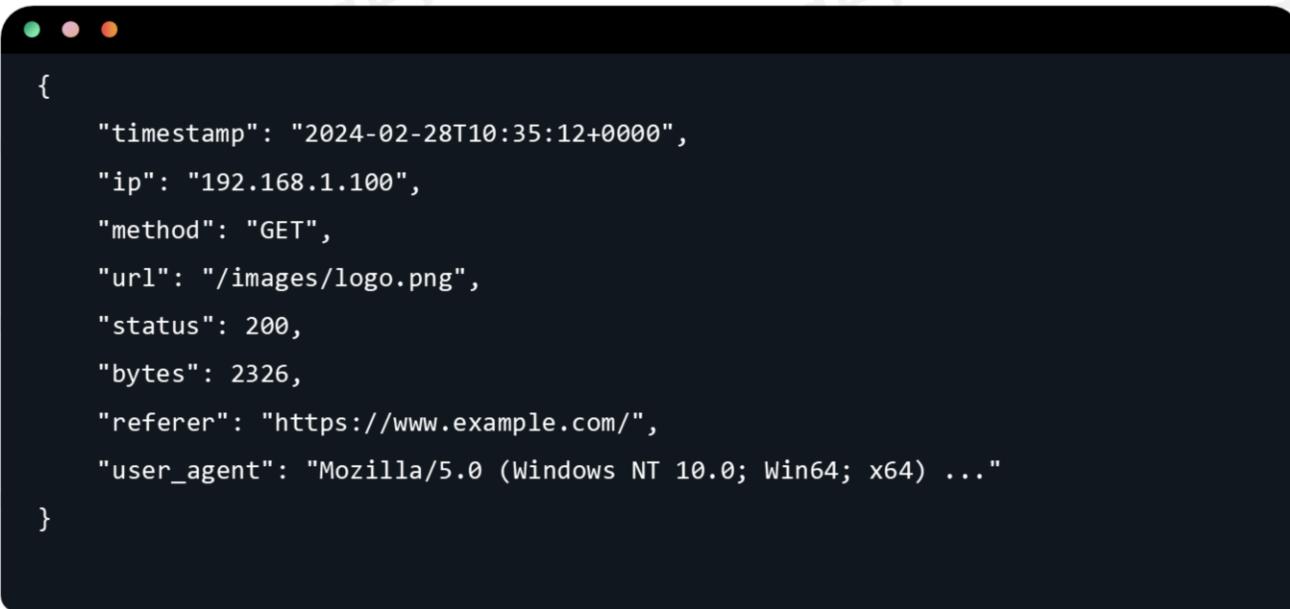
Schema-Less



© Copyright KodeKloud

- **JSON Format:** Documents are stored in JSON format, which is a lightweight data interchange format.
- **Indexing:** Documents are indexed within an Elasticsearch index, which is a collection of similar types of documents.
- **Node:** Documents are stored on nodes, which are individual servers that are part of a larger Elasticsearch cluster.
- **Sharding and Replicas:** Data is split across multiple nodes using sharding, and replicas are copies of an index's shards to ensure data availability and redundancy.
- **Schema-Less:** Elasticsearch does not require a predefined schema for documents, allowing for flexible data structures and dynamic mapping

Documents in Elasticsearch



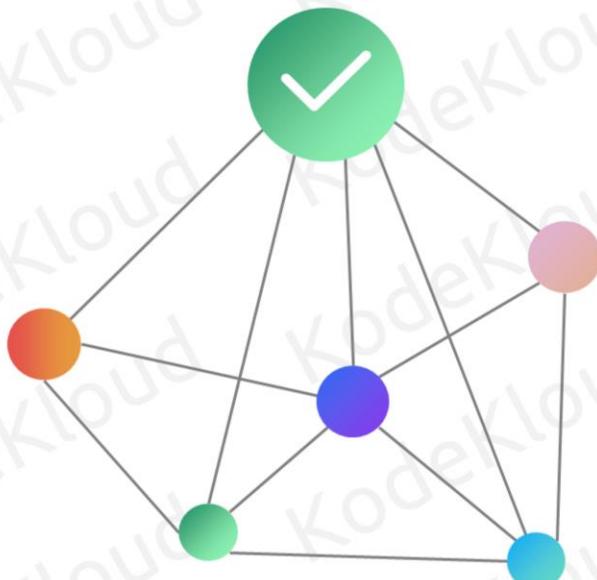
```
{  
    "timestamp": "2024-02-28T10:35:12+0000",  
    "ip": "192.168.1.100",  
    "method": "GET",  
    "url": "/images/logo.png",  
    "status": 200,  
    "bytes": 2326,  
    "referer": "https://www.example.com/",  
    "user_agent": "Mozilla/5.0 (Windows NT 10.0; Win64; x64) ..."  
}
```

Entry in a log file

Documents in Elasticsearch

```
{  
    "_index": "webserver_logs",  
    "_type": "_doc",  
    "_id": "DD46948BFqy1faFsINF8",  
    "_score": 0.2876821,  
    "_source": {  
        "timestamp": "2024-02-28T10:35:12+0000",  
        "ip": "192.168.1.100",  
        "method": "GET",  
        "url": "/images/logo.png",  
        "status": 200,  
        "bytes": 2326,  
        "referer": "https://www.example.com/",  
        "user_agent": "Mozilla/5.0 (Windows NT 10.0; Win64; x64) ..."  
    }  
}
```

Inverted Index in Elasticsearch



Elasticsearch uses an **inverted index** for rapid full-text searches.

- An inverted index is the core data structure used by Elasticsearch to enable fast full-text searches.

Inverted Index in Elasticsearch

Book's Index

Term	Page Numbers
Python	24, 56, 78
Java	12, 45, 67, 89
Database	3, 34, 56, 78, 90
Algorithm	1, 23, 45, 67, 89

Inverted Index

Term	Document IDs
Python	101, 202, 303
Java	303, 404, 505
Database	404, 505, 606
Algorithm	707, 808, 909

Similar to a book's index, Inverted Index maps terms to the documents where they appear.

- It is similar to the index at the end of a book that maps terms to the documents they appear in.

Inverted Index in Elasticsearch

Let's say you have three documents in Elasticsearch:

Document 1: "The quick brown fox jumps over the lazy dog."

Document 2: "The dog chased the cat."

Document 3: "The fox is cunning."



Inverted Index structure

Term	Document IDs
the	1, 2, 3
quick	1
brown	1
fox	1, 3
jumps	1
over	1
lazy	1
dog	1, 2
chased	2
cat	2
is	3
cunning	3

Inverted Index in Elasticsearch

How It Works

01

Tokenization

Elasticsearch divides each document into terms (words).

02

Normalization

Terms are converted to lowercase and processed for consistency.

03

Stop Words Removal

Common words like "the" are removed to improve search relevance.

04

Index Creation

Elasticsearch builds an inverted index mapping terms to documents.

© Copyright KodeKloud

How It Works

Tokenization: Elasticsearch breaks down each document into individual terms (words).

Normalization: It converts terms to lowercase and performs other operations to ensure consistency.

Stop Words Removal: Common words like "the" are often removed as they don't add much value to searches.

Index Creation: The inverted index is built, associating each term with the documents it appears in.

Inverted Index in Elasticsearch



Conduct a search for "fox dog"

It returns document 1 as the result.



Inverted Index structure

Term	Document IDs
the	1, 2, 3
quick	1
brown	1
fox	1, 3
jumps	1
over	1
lazy	1
dog	1, 2
chased	2
cat	2
is	3
cunning	3

© Copyright KodeKloud

Searching with an Inverted Index

When you search for "fox dog", Elasticsearch consults the inverted index:

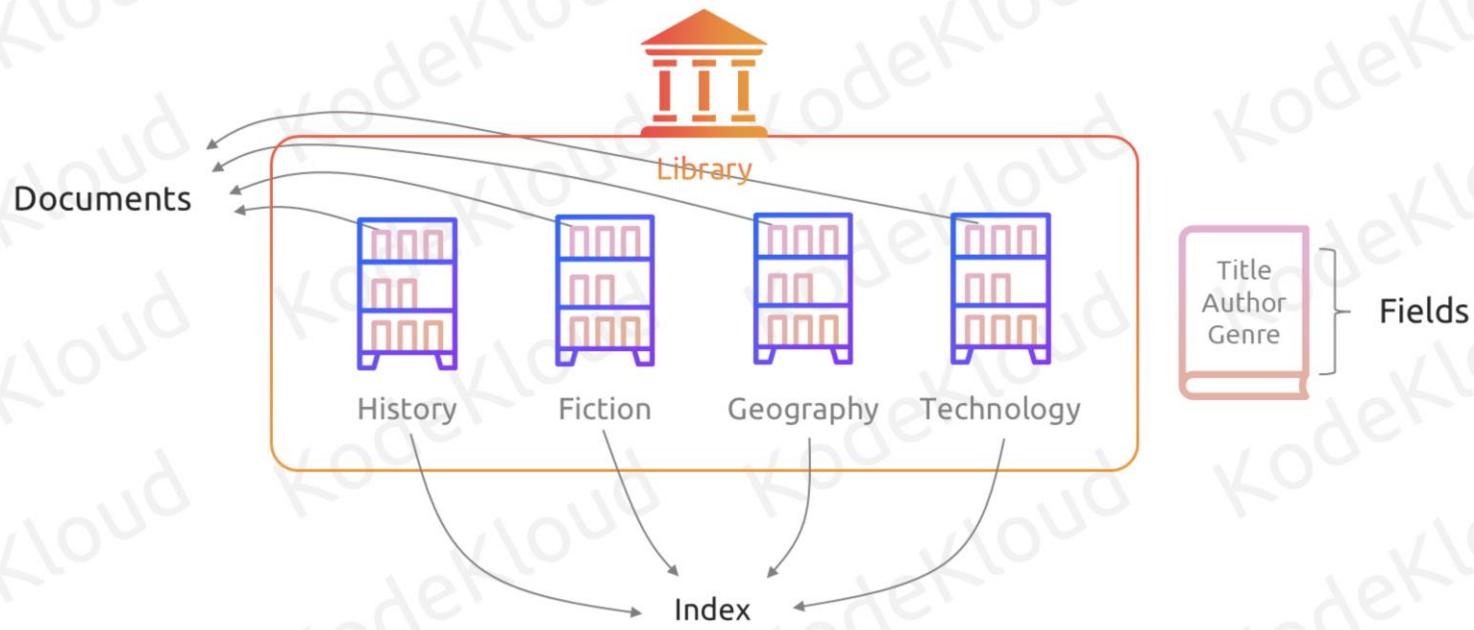
It looks up the document IDs associated with "fox" (1, 3).

It looks up the document IDs associated with "dog" (1, 2).

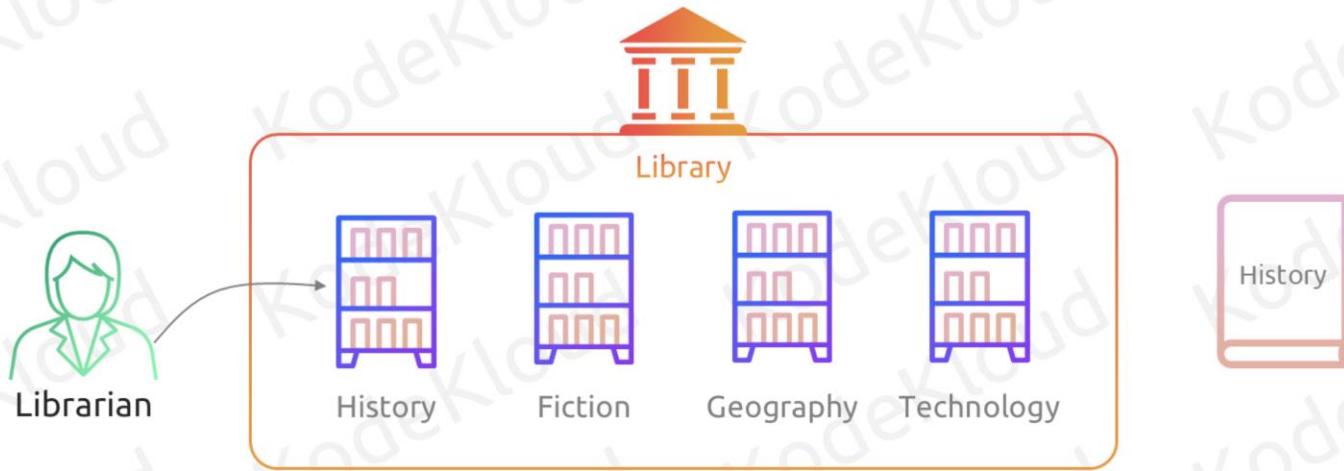
It finds the **intersection** of these two sets of document IDs (1).

It returns document 1 as the result.

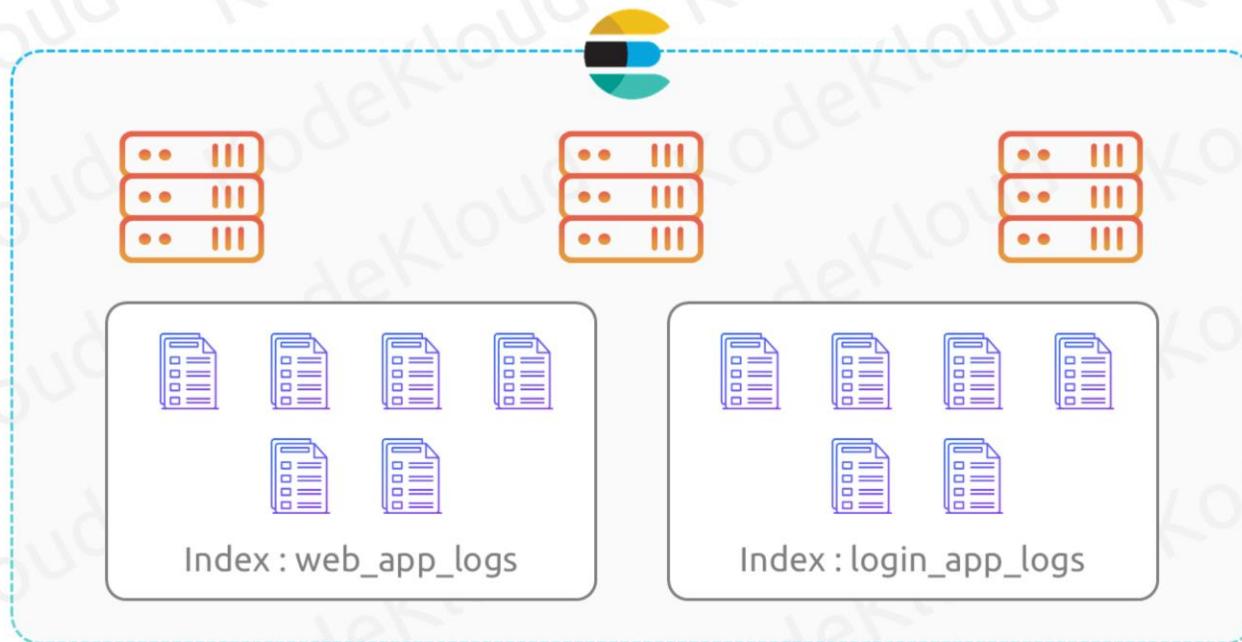
Index in Elasticsearch



Index in Elasticsearch



Index in Elasticsearch





Index in Elasticsearch

01

Logical Container

02

Documents and Fields

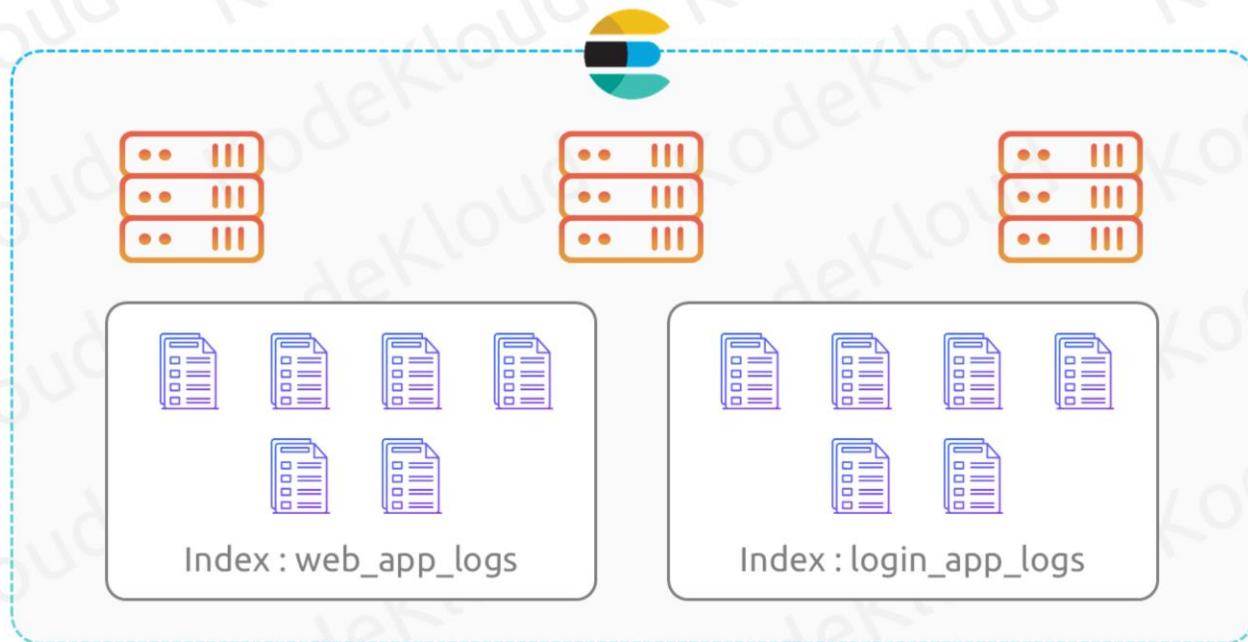
03

Optimized for Search

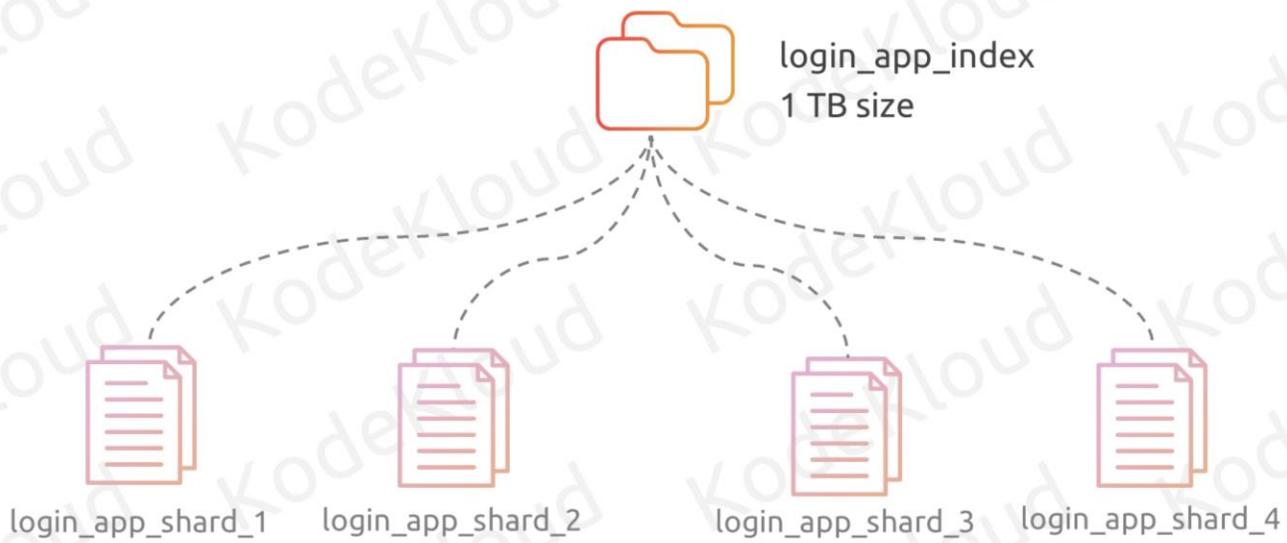
04

Schema Flexibility

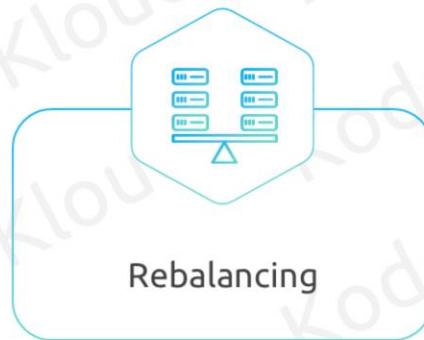
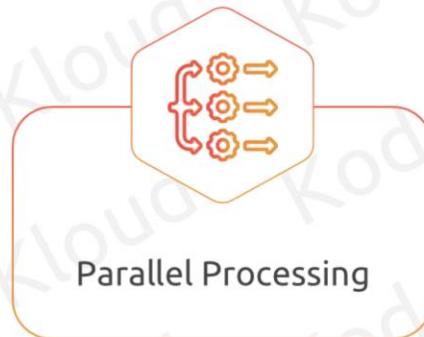
Shards and Replicas in Elasticsearch



Shards and Replicas in Elasticsearch



Shards in Elasticsearch



© Copyright KodeKloud

Data Distribution: Shards are the fundamental building blocks of an Elasticsearch index. They enable horizontal scaling by dividing the index's data into smaller, manageable chunks.

Parallel Processing: Each shard is a self-contained Lucene index, capable of independently performing search and indexing operations. This parallelism boosts performance and allows Elasticsearch to handle large datasets and high query volumes.

Fault Tolerance: Shards can be replicated across multiple nodes for redundancy. If a node fails, the data in its shards can be

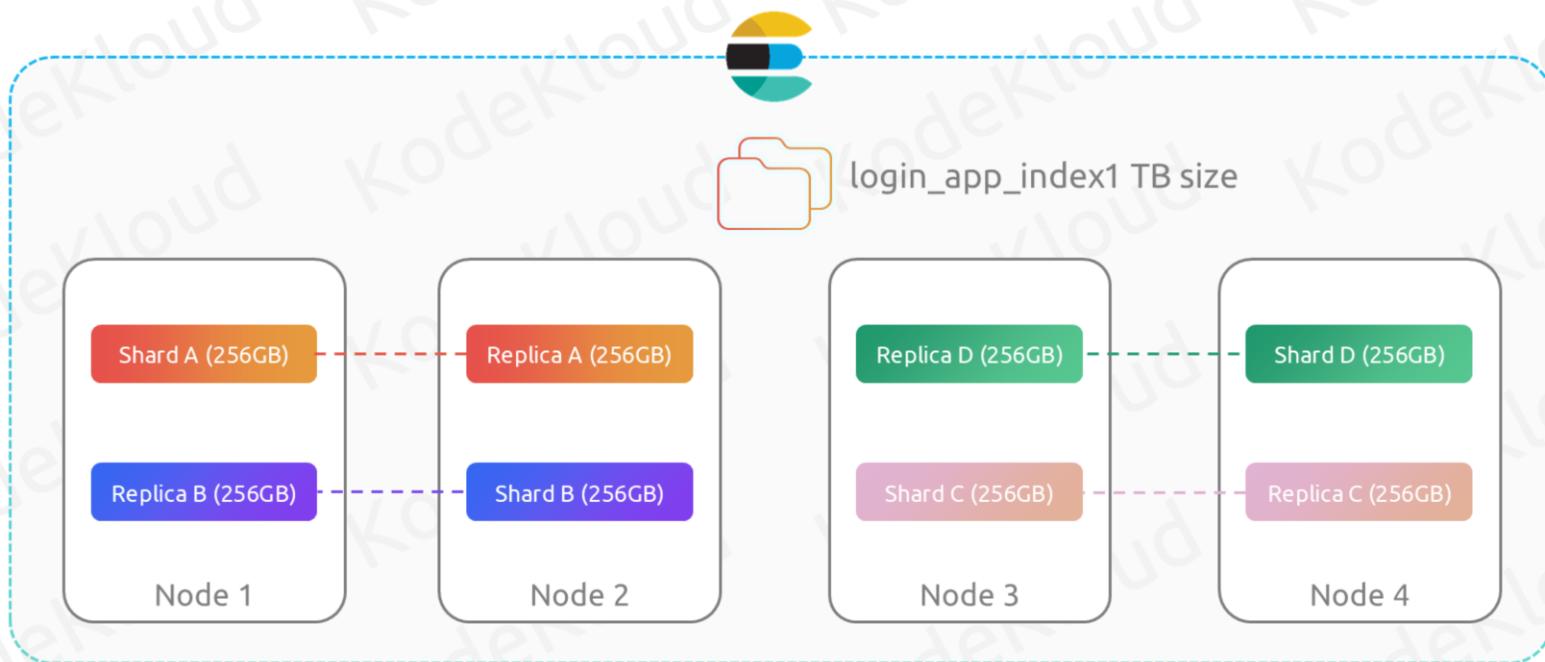
recovered from replicas, ensuring high availability.

Scalability and Performance: By adding more nodes to the cluster, you can increase the number of shards, effectively distributing data and workload, and improving both scalability and performance.

Rebalancing: Elasticsearch automatically rebalances shards across nodes when the cluster configuration changes (e.g., adding or removing nodes). This ensures even data distribution and optimal resource utilization.

Shard Sizing: It's important to choose an appropriate shard size based on your data volume and growth expectations. Too few large shards might hinder performance and scalability, while too many small shards could lead to excessive overhead.

Shards and Replicas in Elasticsearch



Replicas in Elasticsearch



© Copyright KodeKloud

Data Redundancy: Replicas create copies of primary shards, ensuring data availability in case of failures. One-liner: "Replicas: Your data's safety net in Elasticsearch."

High Availability: Replicas allow Elasticsearch to automatically promote a replica to primary if the original fails. One-liner: "Replicas: The unsung heroes of Elasticsearch uptime."

Increased Read Capacity: Replicas can serve search requests, improving performance under heavy read loads. One-liner:

"Replicas: Supercharging your Elasticsearch search speed."

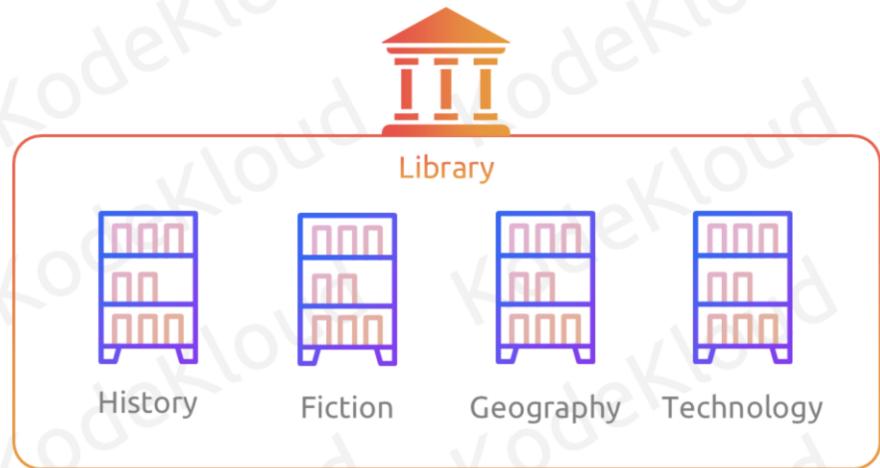
Configurable Per Index: You can customize the number of replicas for each index to match its specific needs. One-liner: "Replicas: Tailor-made data protection for each Elasticsearch index."

Automatic Shard Allocation: Elasticsearch intelligently manages replica placement across nodes for optimal resilience. One-liner: "Replicas: Elasticsearch's smart data distribution strategy."

Shard Recovery: If a node fails, replicas help to quickly restore data to new shards on healthy nodes. One-liner: "Replicas: The first responders of Elasticsearch data recovery."

Zero Downtime Scaling: Adding replicas allows you to scale read capacity without interrupting ongoing operations. One-liner: "Replicas: Your ticket to seamless Elasticsearch expansion."

Mapping in Elasticsearch



Mapping in Elasticsearch

Mapping is the process of defining how a document and the fields it contains are stored and indexed.

Dynamic Mapping

Explicit Mapping

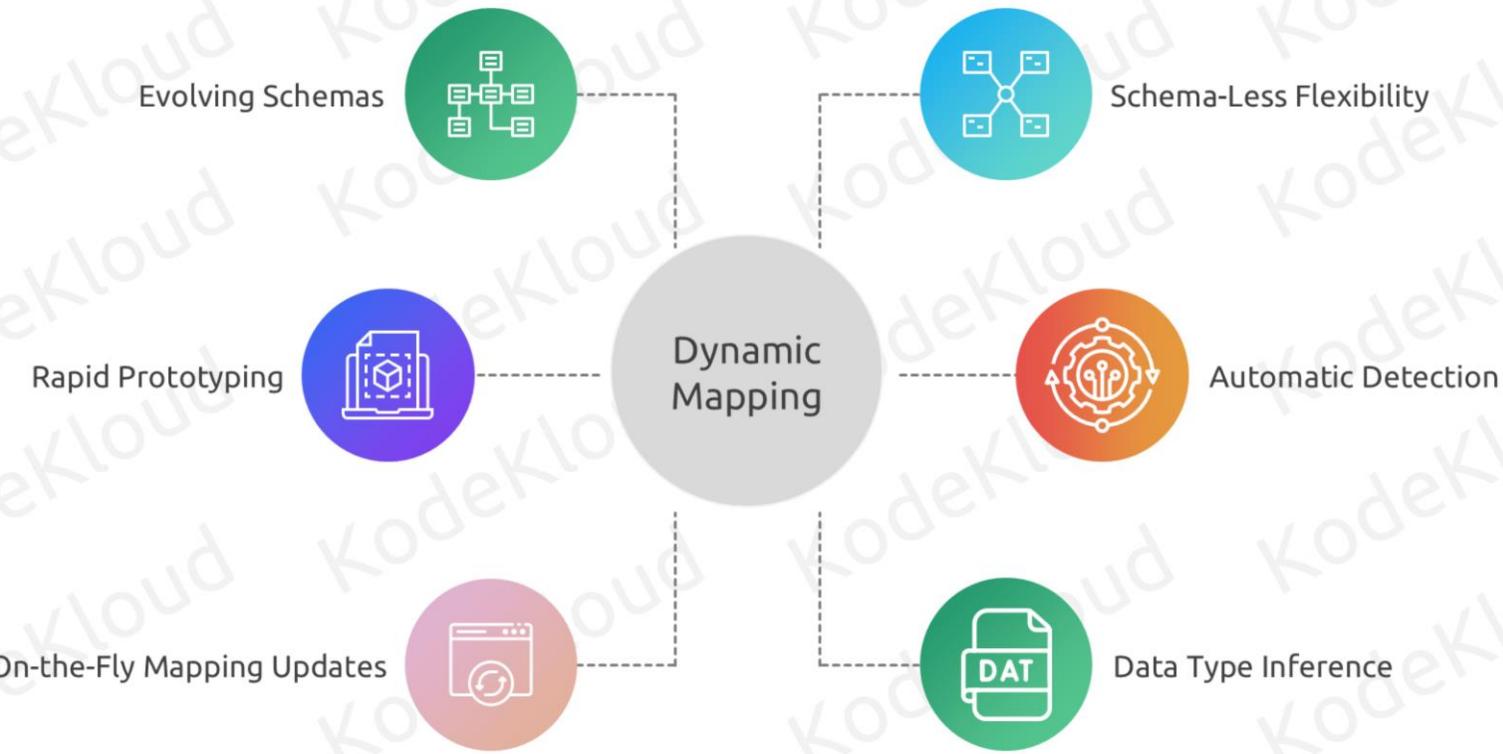
Dynamic Mapping in Elasticsearch

```
{  
    "product_id": "12345",  
    "name": "Awesome Sneakers",  
    "description": "The most comfortable shoes ever!",  
    "price": 99.99,  
    "in_stock": true,  
    "release_date": "2023-12-15"  
}
```



```
{  
    "mappings": {  
        "_doc": {  
            "properties": {  
                "product_id": { "type": "long" },  
                "name": { "type": "text" },  
                "description": { "type": "text" },  
                "price": { "type": "float" },  
                "in_stock": { "type": "boolean" },  
                "release_date": { "type": "date" }  
            }  
        }  
    }  
}
```

Dynamic Mapping in Elasticsearch



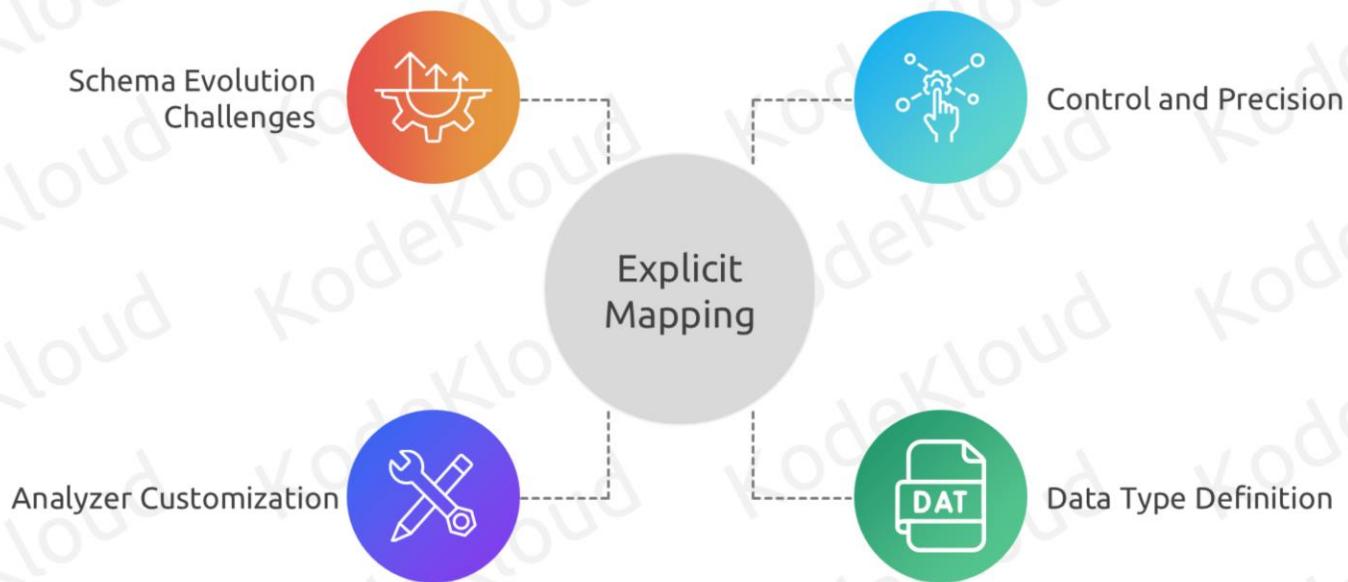
Explicit Mapping in Elasticsearch

```
PUT /product_index
{
  "mappings": {
    "_doc": {
      "properties": {
        "product_id": { "type": "keyword" },
        "name": {
          "type": "text",
          "analyzer": "english"
        },
        "description": { "type": "text" },
        "price": { "type": "float" },
        "in_stock": { "type": "boolean" },
        "release_date": {
          "type": "date",
          "format": "yyyy-MM-dd"
        }
      }
    }
  }
}
```

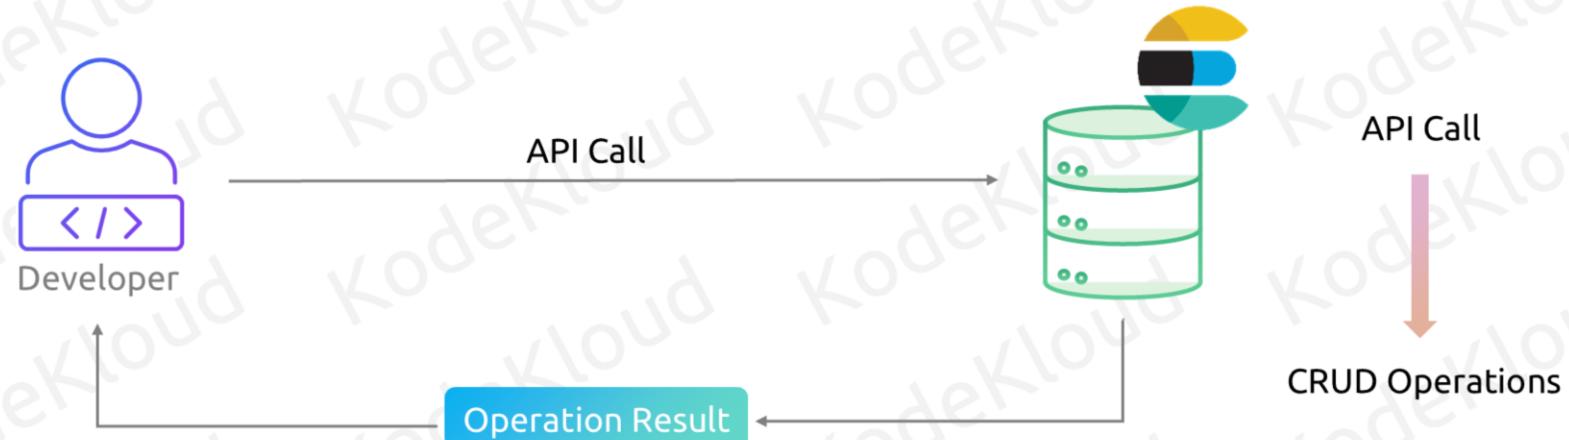


```
{
  "product_id": "12345",
  "name": "Awesome Sneakers",
  "description": "The most comfortable shoes ever!",
  "price": 99.99,
  "in_stock": true,
  "release_date": "2023-12-15"
}
```

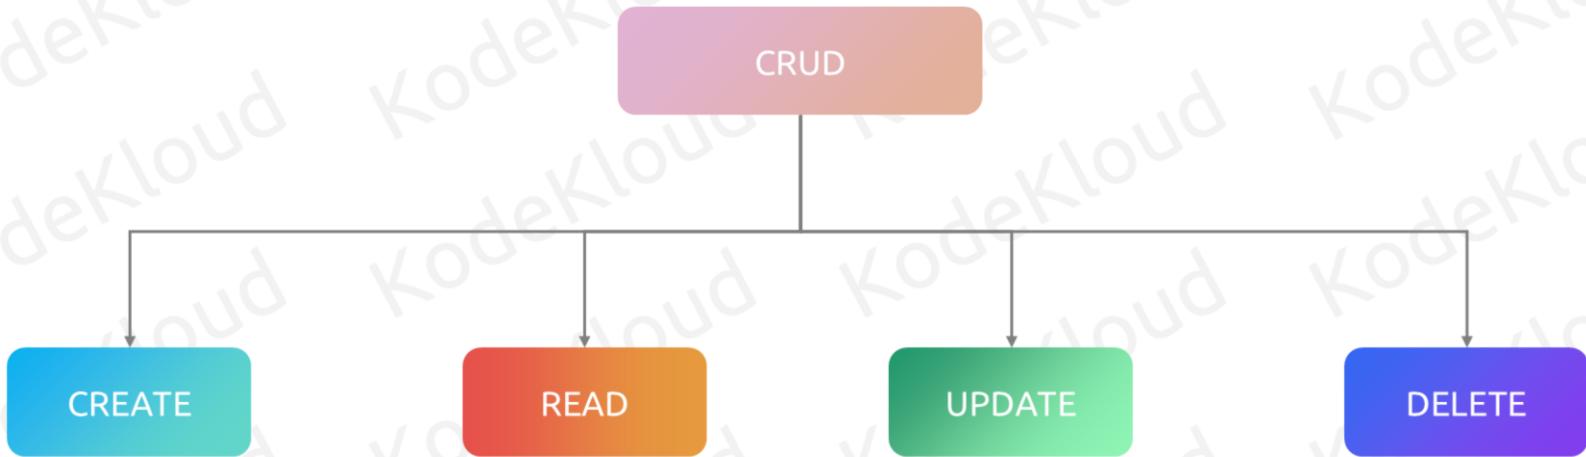
Mapping in Elasticsearch



CRUD Operations on Elasticsearch



CRUD Operations on Elasticsearch



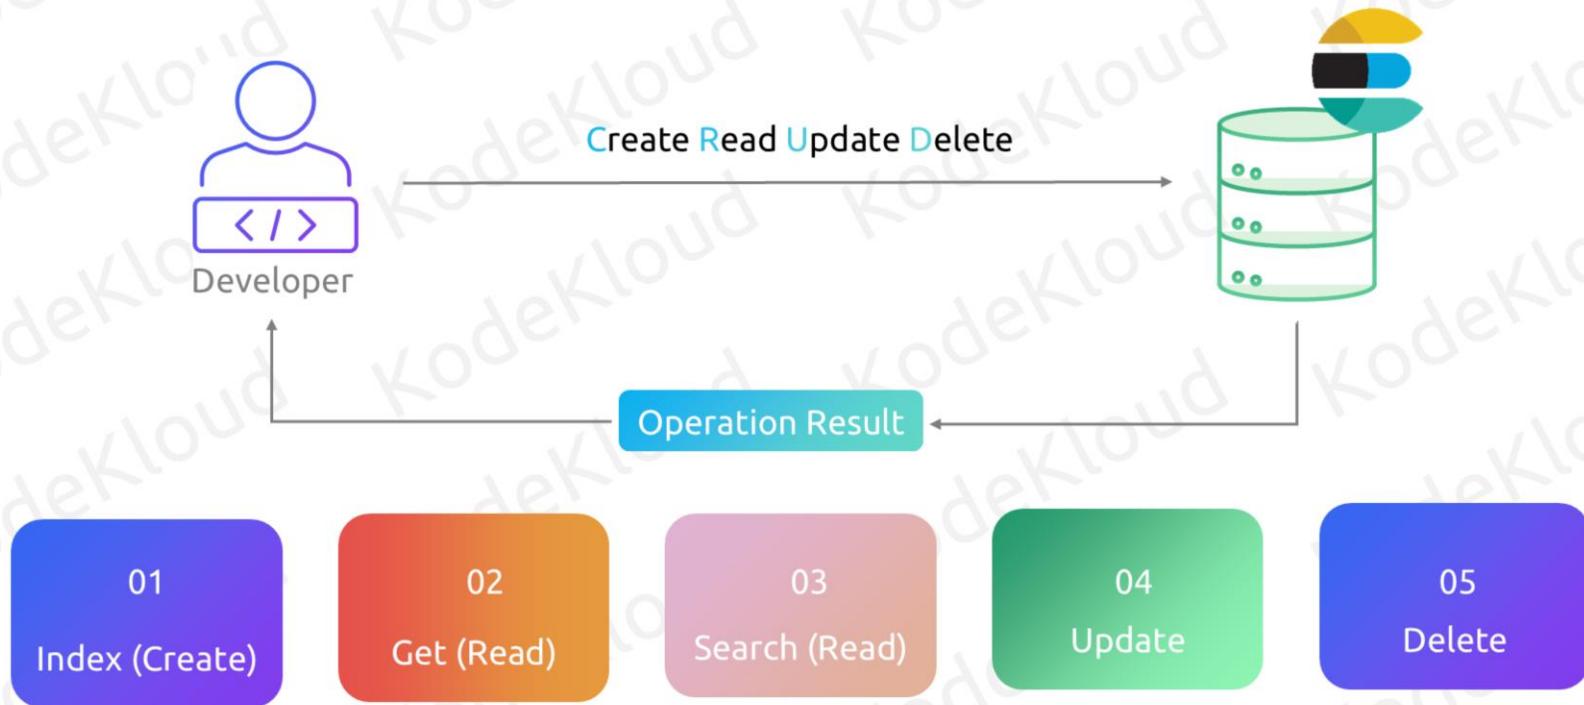
© Copyright KodeKloud

What CRUD Stands For

CRUD is an acronym that describes the four basic operations you can perform on data in a database or search engine like Elasticsearch:

- Create
- Read
- Update
- Delete

CRUD Operations on Elasticsearch



© Copyright KodeKloud

Index (Create): Adds new JSON documents to a specific index, like adding books to a library shelf. Each document gets a unique ID for later retrieval.

Get (Read): Retrieves a specific document by its unique ID from the relevant index, similar to finding a book in a library by its call number.

Search (Read): Finds documents based on their content using flexible queries. This is like searching the library catalog by keywords or author.

Update: Modifies existing documents, either by updating specific fields or by changing multiple documents that match a

search query.

Delete: Removes documents from the index either by their unique ID or by matching a query. This is equivalent to removing a book from the library.

Demo

Elasticsearch CRUD
commands

Demo

Difference between
POST and PUT operations

Demo

Cluster Information –
Elasticsearch CRUD
commands

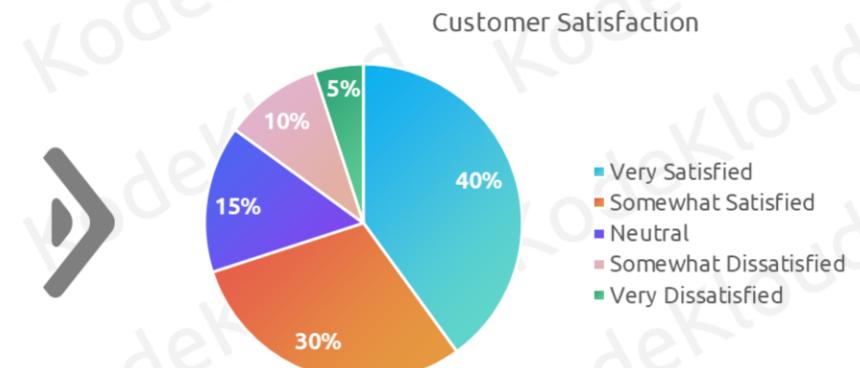
Understanding Kibana

Importance of Data Visualization

Rating	Number of Customers	Percentage
Very Satisfied	80	40%
Somewhat Satisfied	60	30%
Neutral	30	15%
Somewhat Dissatisfied	20	10%
Very Dissatisfied	10	5%
Total	200	100%

Importance of Data Visualization

Rating	Number of Customers	Percentage
Very Satisfied	80	40%
Somewhat Satisfied	60	30%
Neutral	30	15%
Somewhat Dissatisfied	20	10%
Very Dissatisfied	10	5%
Total	200	100%



Importance of Data Visualization

01



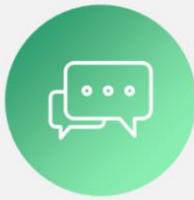
Improved
understanding

02



Faster decision
making

03



Effective
communication

04



Identifying
hidden patterns

05



Data-driven
storytelling

© Copyright KodeKloud

Improved Understanding: Data visualization transforms raw data into intuitive formats like charts, graphs, and maps. This makes complex trends, patterns, and outliers much easier to grasp.

Faster Decision Making: Visualizations provide a quick overview of the data, allowing decision-makers to identify key insights and act upon them swiftly. This is crucial in fast-paced business environments.

Effective Communication: Visuals are a powerful way to communicate findings to both technical and non-technical audiences. They make data more engaging and memorable, aiding in conveying complex information effectively.

Identifying Hidden Patterns: Data visualization can reveal hidden relationships and correlations that might not be apparent

in raw data. This can lead to new discoveries and opportunities.

Data-Driven Storytelling: Visualizations help weave compelling narratives around data, making it easier to communicate the story behind the numbers. This is especially useful for presentations and reports.

Kibana Visualizations and Dashboards

Basic Charts

Time Series Visualizations

Geospatial Visualizations

Table: Organizes data in rows and columns for detailed inspection

Month	Product A Sales	Product B Sales	Website Traffic	Customer Satisfaction (1-10)	Marketing Spend
January	5,000	3,500	20,000	7.5	\$500
February	6,200	4,200	25,000	8	\$650
March	7,800	5,500	32,000	8.2	\$800
April	6,500	4,800	28,000	7.8	\$700
May	8,200	6,000	35,000	8.5	\$850

© Copyright KodeKloud

Basic Charts:

- **Area Chart:** Displays trends over time, with the area between the line and the axis filled.
- **Bar Chart:** Shows comparisons between categories or values.
- **Heatmap:** Visualizes data density across two dimensions using colors.
- **Line Chart:** Tracks changes in values over time.
- **Metric:** Displays a single numeric value.
- **Pie Chart:** Shows proportions of a whole.

- **Scatterplot:** Plots points based on two variables to reveal relationships.
- **Table:** Organizes data in rows and columns for detailed inspection.

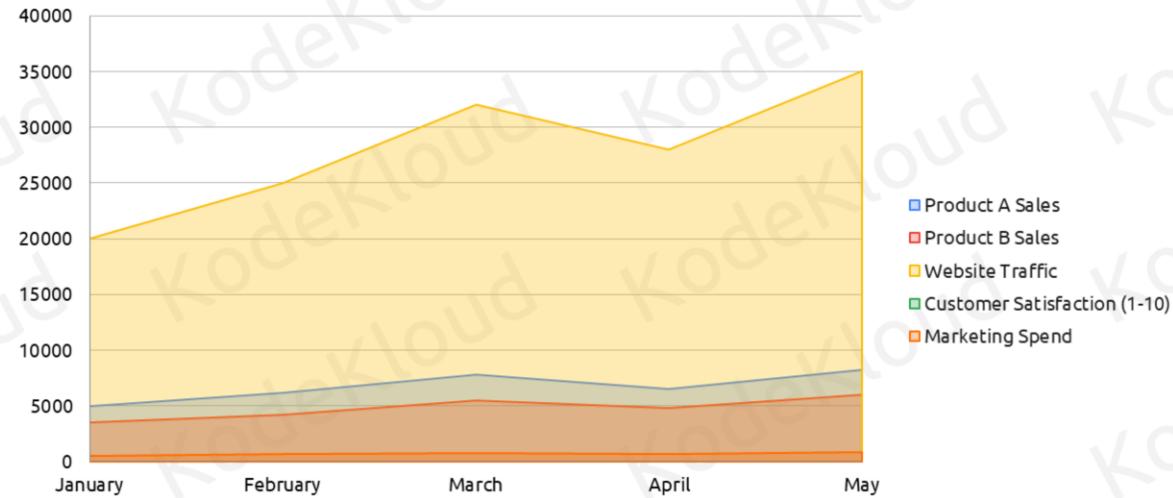
Kibana Visualizations and Dashboards

Basic Charts

Time Series Visualizations

Geospatial Visualizations

Area Chart: Displays trends over time, with the area between the line and the axis filled



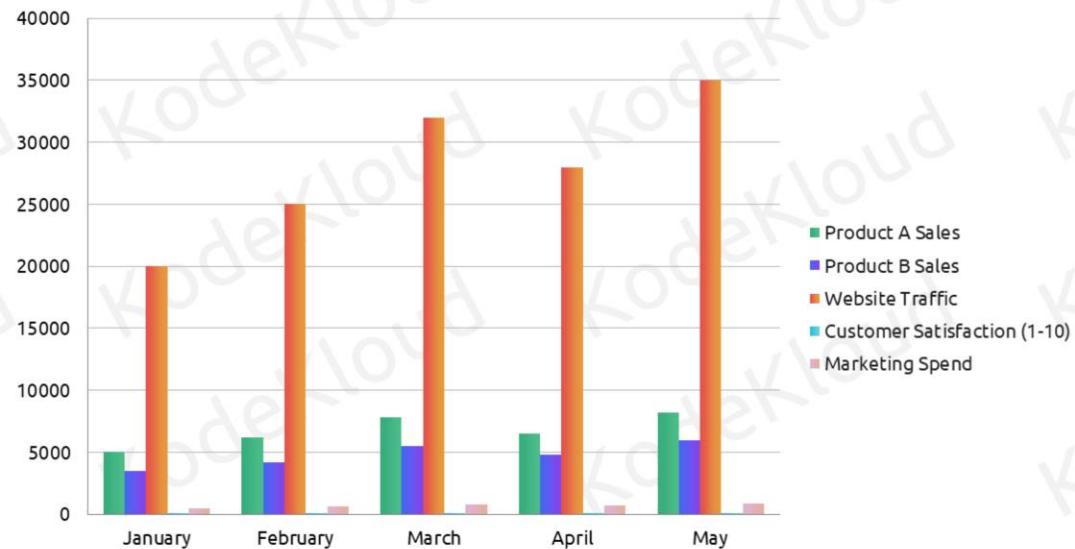
Kibana Visualizations and Dashboards

Basic Charts

Time Series Visualizations

Geospatial Visualizations

Bar Chart: Shows comparisons between categories or values



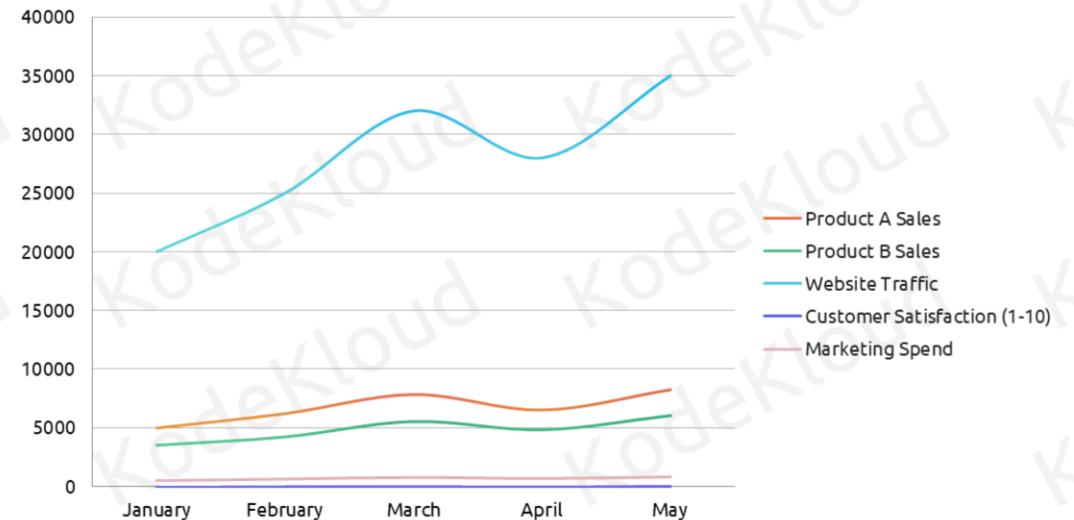
Kibana Visualizations and Dashboards

Basic Charts

Time Series Visualizations

Geospatial Visualizations

Line Chart: Tracks changes in values over time



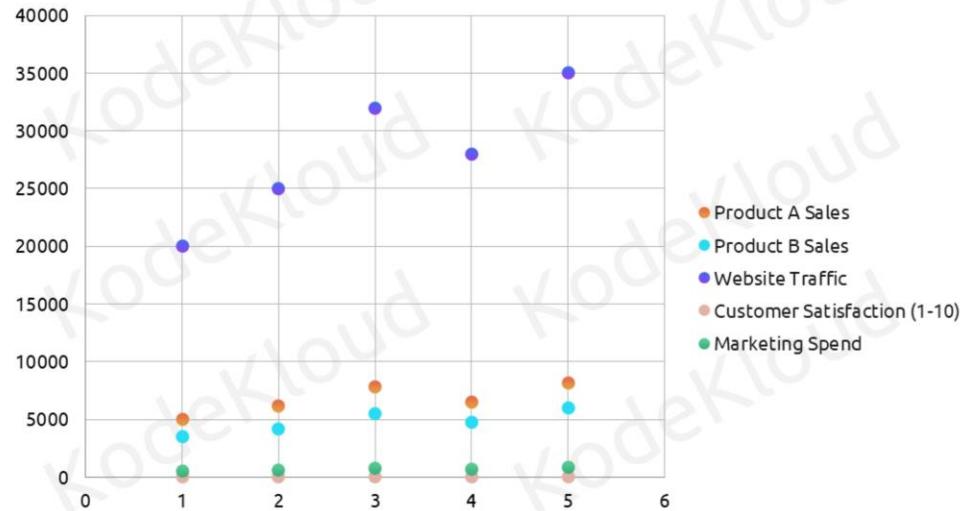
Kibana Visualizations and Dashboards

Basic Charts

Time Series Visualizations

Geospatial Visualizations

Scatterplot: Plots points based on two variables to reveal relationships: Tracks changes in values over time



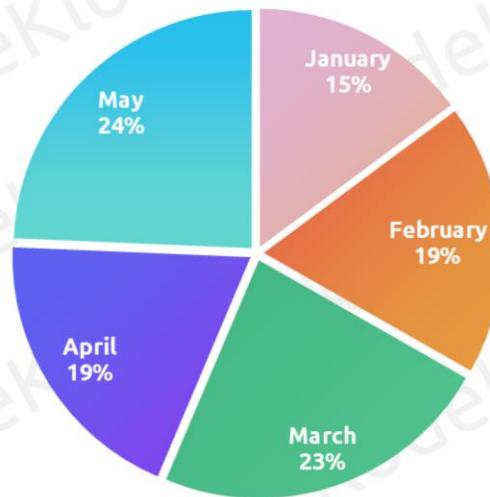
Kibana Visualizations and Dashboards

Basic Charts

Time Series Visualizations

Geospatial Visualizations

Pie Chart: Shows proportions of a whole



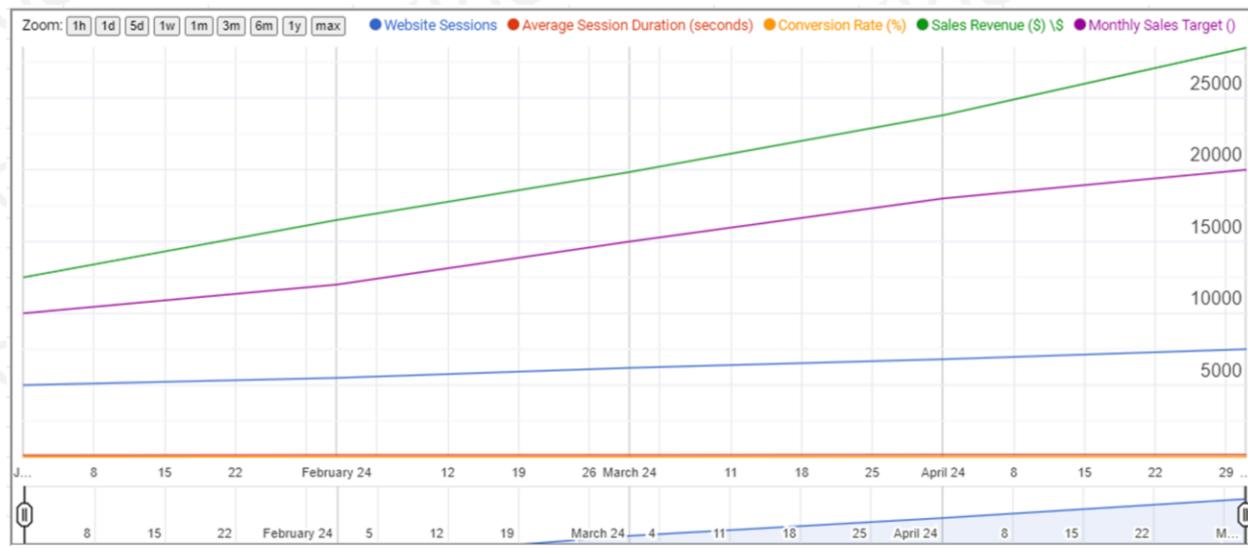
Kibana Visualizations and Dashboards

Basic Charts

Time Series Visualizations

Geospatial Visualizations

Time Series Visual Builder (TSVB): Creates sophisticated visualizations with multiple metrics and complex aggregations over time; ideal for analyzing trends, seasonality, and anomalies



© Copyright KodeKloud

Time Series Visualizations:

- **Time Series Visual Builder (TSVB):** Creates sophisticated visualizations with multiple metrics and complex aggregations over time. Ideal for analyzing trends, seasonality, and anomalies.
- **Gauge:** Displays a single value on a circular scale, often used for performance metrics.
- **Goal:** Tracks progress towards a target value.

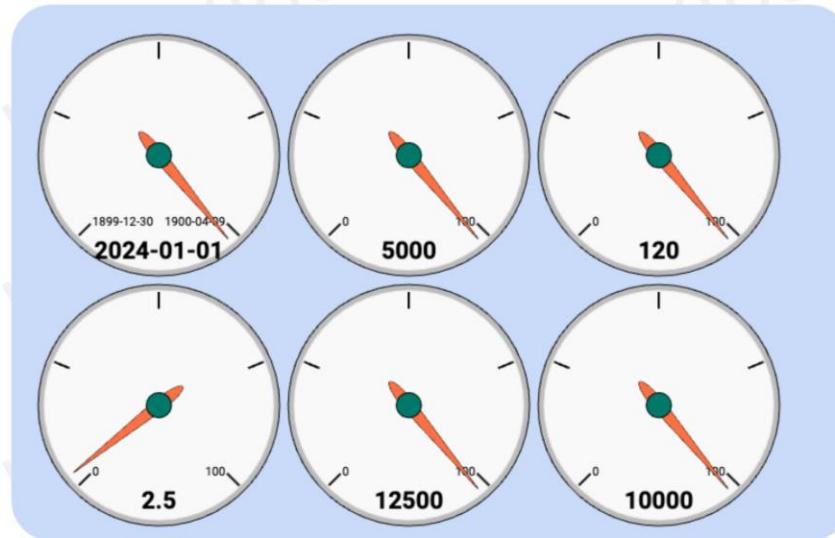
Kibana Visualizations and Dashboards

Basic Charts

Time Series Visualizations

Geospatial Visualizations

Gauge: Displays a single value on a circular scale, often used for performance metrics



© Copyright KodeKloud

Time Series Visualizations:

- **Time Series Visual Builder (TSVB):** Creates sophisticated visualizations with multiple metrics and complex aggregations over time. Ideal for analyzing trends, seasonality, and anomalies.
- **Gauge:** Displays a single value on a circular scale, often used for performance metrics.
- **Goal:** Tracks progress towards a target value.

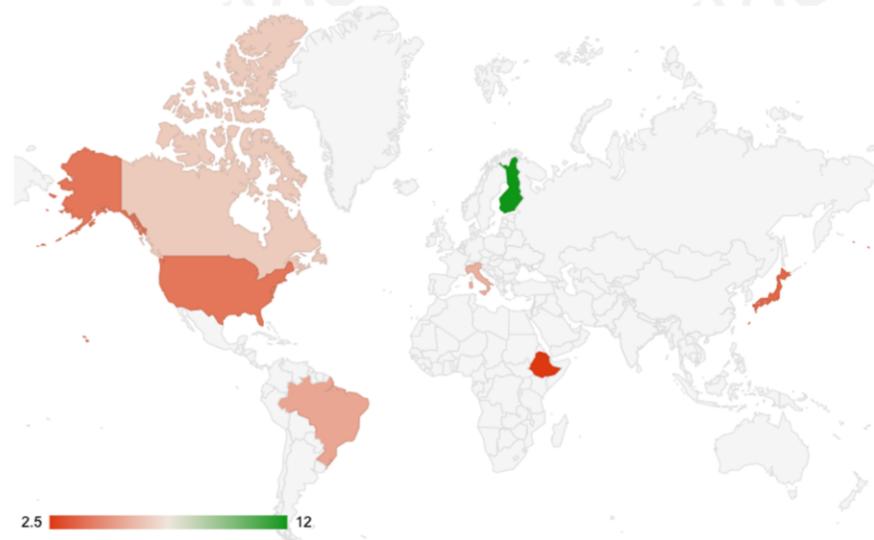
Kibana Visualizations and Dashboards

Basic Charts

Time Series Visualizations

Geospatial Visualizations

Region Map: Colors geographical regions based on data values

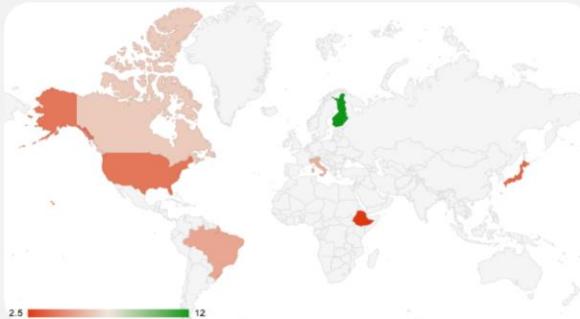
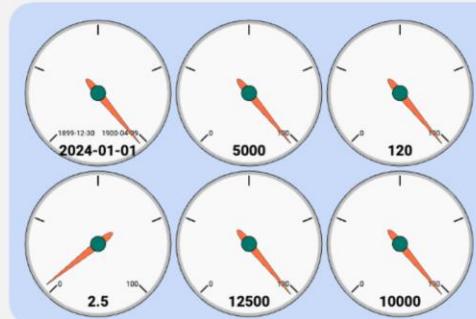
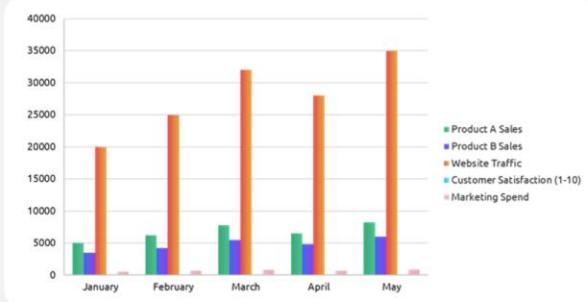
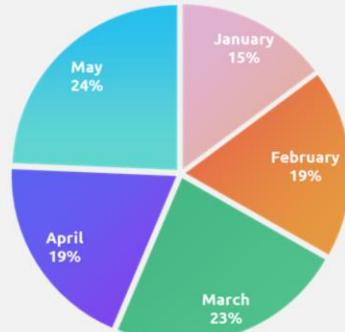


© Copyright KodeKloud

Geospatial Visualizations:

- Region Map: Colors geographical regions based on data values.

Kibana Visualizations and Dashboards



Demo

Building our first
dashboard using Kibana

Kibana Query Language (KQL) – Mastering Syntax and Advanced Querying

Querying Elasticsearch Data



- KQL is integrated into Kibana for querying Elasticsearch data.
- KQL queries are translated into Elasticsearch Query DSL requests.
- Elasticsearch processes these requests and returns matching data to Kibana.

© Copyright KodeKloud

- **Querying Elasticsearch Data**
KQL is a query language built into Kibana, which is a visualization and exploration tool for data stored in Elasticsearch.
- When a user enters a KQL query in Kibana, it translates this query into an Elasticsearch Query DSL (Domain Specific Language) request. This request is then sent to the Elasticsearch cluster.
- Elasticsearch processes the request, searches through the indexed data, and returns the matching documents to Kibana.

Kibana Query Language (KQL) – Mastering Syntax and Advanced Querying

Enter a Query in Kibana:

```
status: "200" AND  
extension: "php"
```



Translation to Elasticsearch Query DSL:

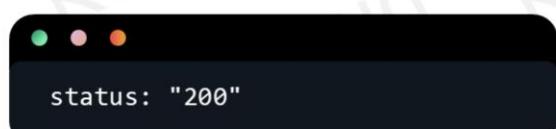
```
{  
  "query": {  
    "bool": {  
      "must": [  
        { "match": { "status": "200" }},  
        { "match": { "extension": "php" }}  
      ]  
    }  
  }  
}
```

KQL – Basic Syntax

- **Field Queries**

Syntax:

`field_name: "value"`



```
status: "200"
```

© Copyright KodeKloud

1. **Field Queries**:

- Syntax: `field_name: "value"`

- Example:

```
```kql
```

```
status: "200"
```

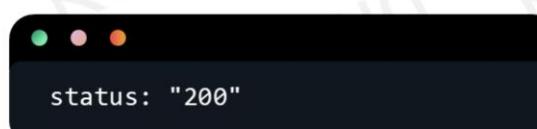
```
``` This query will return documents where the `status` field has the value `200`.
```

KQL – Basic Syntax

- **Field Queries**

Syntax:

`field_name: "value"`

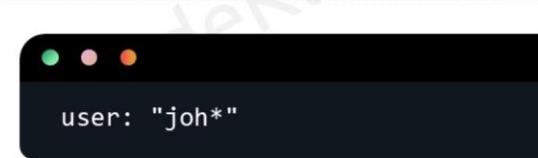


```
status: "200"
```

- **Wildcard Queries**

Syntax:

`field_name: "val*"`



```
user: "joh*"
```

© Copyright KodeKloud

2. **Wildcard Queries**:

- Syntax: `field_name: "val*"`

- Example:

```
```kql
```

```
user: "joh*"
```

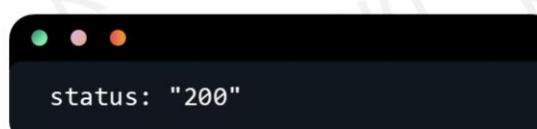
```
``` This query will return documents where the `user` field starts with "joh", such as "john" or "johnny".
```

KQL – Basic Syntax

- **Field Queries**

Syntax:

`field_name: "value"`



```
status: "200"
```

- **Wildcard Queries**

Syntax:

`Field_name: "val*"`

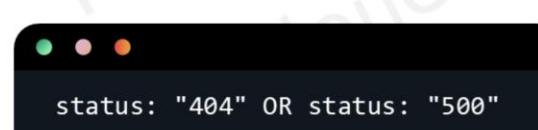


```
user: "joh*"
```

- **Logical Operators**

Syntax:

`condition1 AND/OR/NOT condition2`



```
status: "404" OR status: "500"
```

© Copyright KodeKloud

3. **Logical Operators**:

- Syntax: `condition1 AND/OR/NOT condition2`

- Example:

```
```kql
```

```
status: "404" OR status: "500"
```

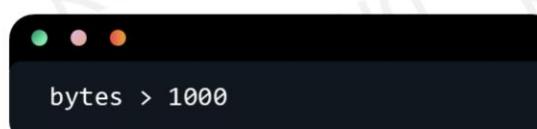
```
``` This query will return documents where the `status` field is either `404` or `500`.
```

KQL – Basic Syntax

• Regular Expressions •

Syntax:

`field_name > value or
field_name < value`



```
bytes > 1000
```

© Copyright KodeKloud

4. **Range Queries**:

- Syntax: `field_name > value` or `field_name < value`

- Example:

```
```kql
```

```
bytes > 1000
```

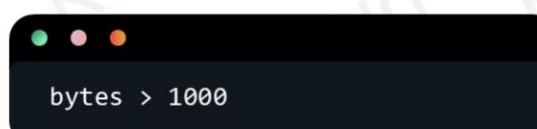
``` This query will return documents where the `bytes` field has a value greater than `1000`.

KQL – Basic Syntax

- **Regular Expressions** •

Syntax:

`field_name > value or
field_name < value`

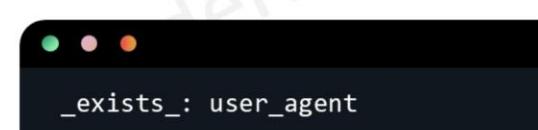


```
bytes > 1000
```

- **Existence Queries** •

Syntax:

`_exists_: field_name`



```
_exists_: user_agent
```

5. **Existence Queries**:

- Syntax: `'_exists_: field_name'`

- Example:

```
```kql
```

```
exists: user_agent
```

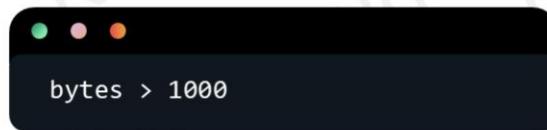
```
``` This query will return documents that contain the `user_agent` field.
```

KQL – Basic Syntax

- **Regular Expressions** •

Syntax:

`field_name > value or
field_name < value`

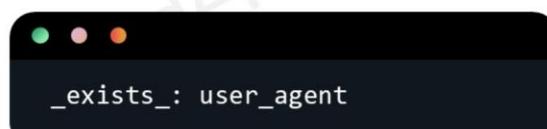


```
bytes > 1000
```

- **Existence Queries** •

Syntax:

`_exists_: field_name`



```
_exists_: user_agent
```

© Copyright KodeKloud

Examples Combined in One Query

Combining different syntax elements, here's a more complex example:

```
```kql
```

```
(status: "200" AND extension: "php") OR (bytes > 1000 AND _exists_: user_agent)```
```

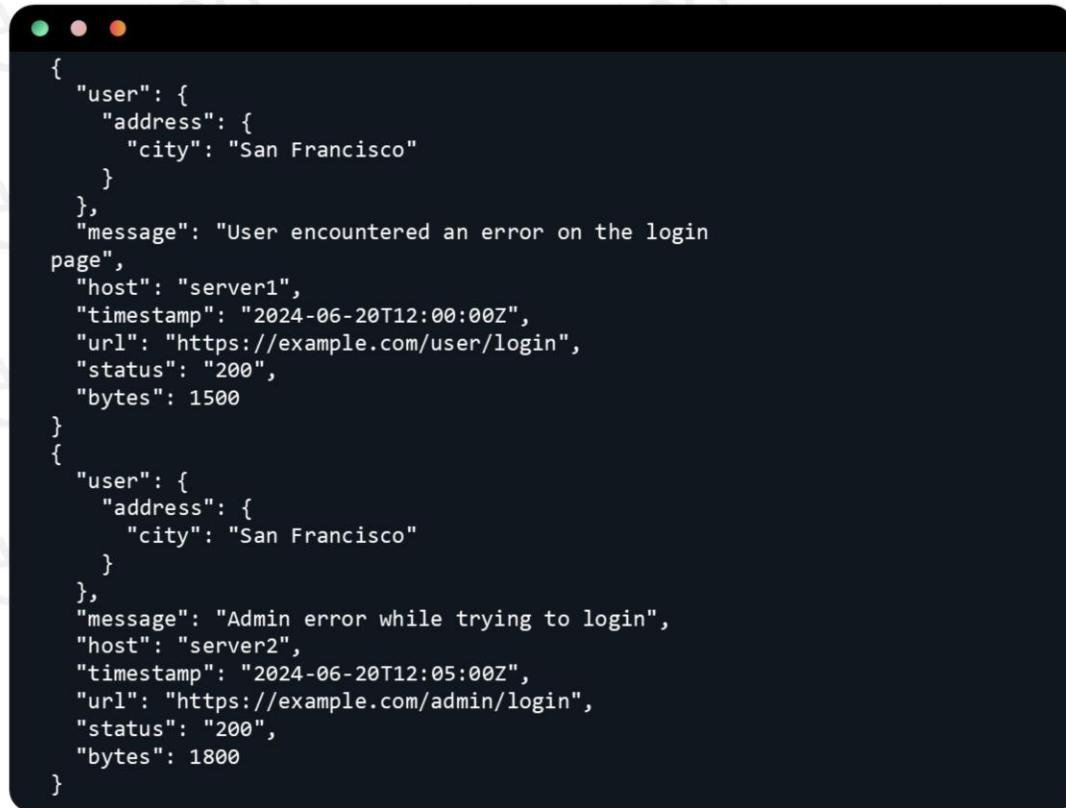
This query will return documents that either have a `status` of `200` and an `extension` of "php", or have more than `1000` bytes and contain the `user\_agent` field.

# KQL – Basic Syntax

Combined in One Query

```
(status: "200" AND extension: "php") OR
(bytes > 1000 AND _exists_: user_agent)
```

# KQL – Advanced Capabilities



```
{
 "user": {
 "address": {
 "city": "San Francisco"
 }
 },
 "message": "User encountered an error on the login
page",
 "host": "server1",
 "timestamp": "2024-06-20T12:00:00Z",
 "url": "https://example.com/user/login",
 "status": "200",
 "bytes": 1500
}

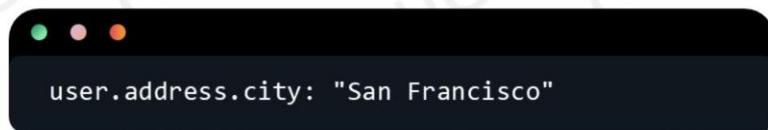
{
 "user": {
 "address": {
 "city": "San Francisco"
 }
 },
 "message": "Admin error while trying to login",
 "host": "server2",
 "timestamp": "2024-06-20T12:05:00Z",
 "url": "https://example.com/admin/login",
 "status": "200",
 "bytes": 1800
}
```

# KQL – Advanced Capabilities

- **Nested Fields** •

Syntax:

`parent_field.child_field: "value"`

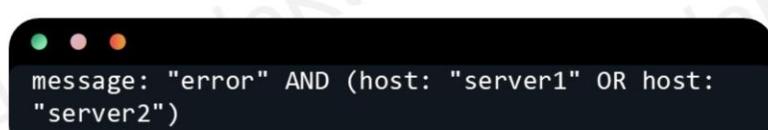


```
user.address.city: "San Francisco"
```

- **Lucene Syntax Support** •

Syntax:

`Field_name: "value" AND (field_name2: "value2"  
OR field_name3: "value3")`



```
message: "error" AND (host: "server1" OR host:
"server2")
```

- **Proximity Searches** •

Syntax:

`"term1 term2"~number`



```
"quick brown"~2
```

# KQL – Advanced Capabilities

## Regular Expressions

Syntax:

field\_name: /regex/



```
url: ./.*login.*/
```

## Boosting

Syntax:

field\_name: "value"^(number)



```
status: "200"^(2)
```

# KQL – Advanced Capabilities

Combined in One Query

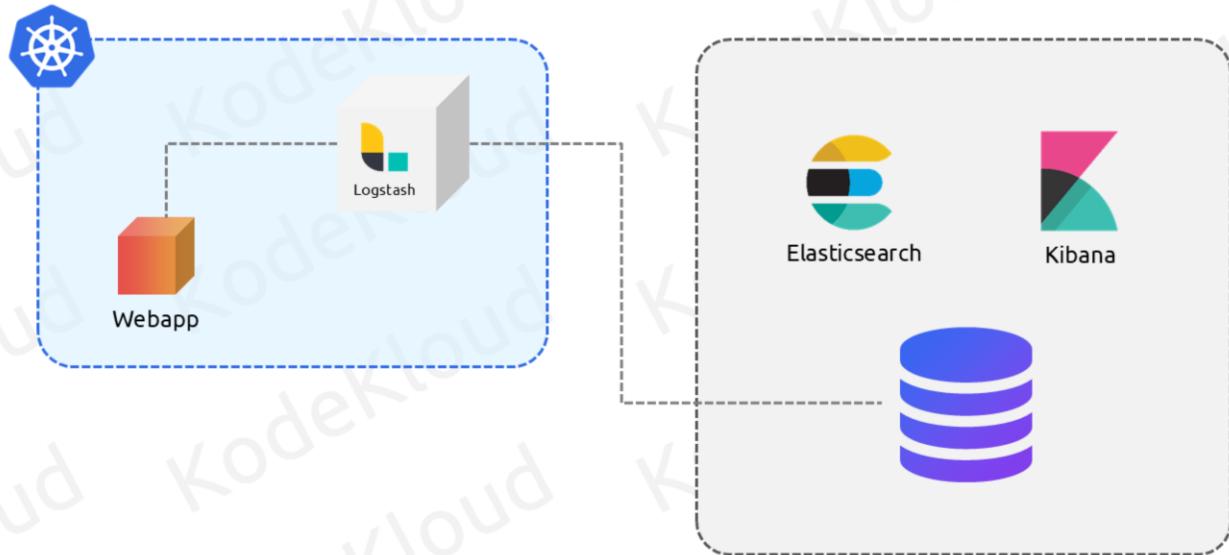
```
(user.address.city: "San Francisco" AND
(message: "error" AND (host: "server1"
OR host: "server2")) AND "quick brown"~2
AND url: /.*login.*/ AND status: "200"~2
```

# Demo

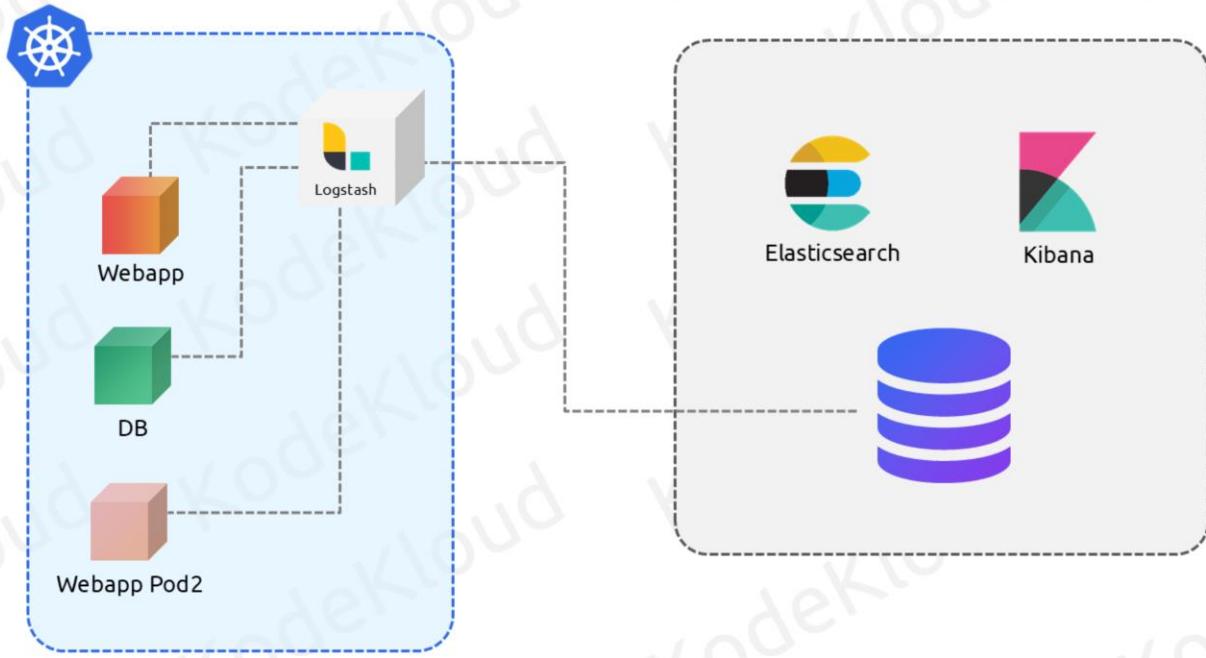
## Kibana Query Language (KQL)

# **FluentBit**

# Logstash – Introduction



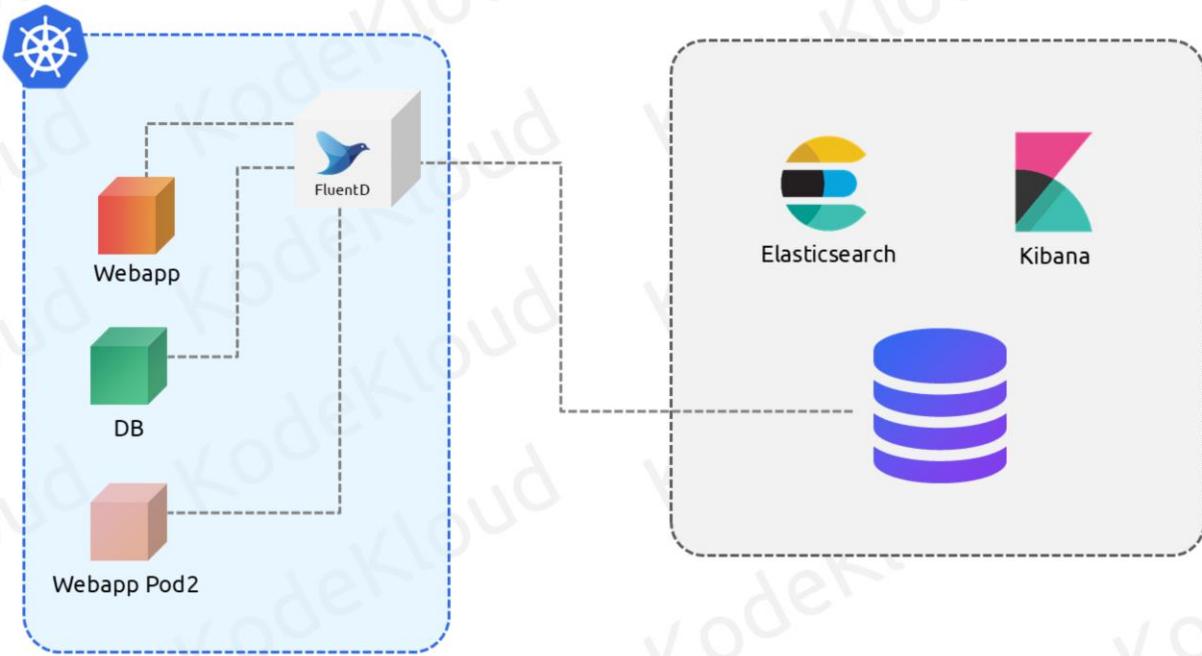
# Logstash – Introduction



# Logstash's Role Within the ELK Stack



# Replacing Logstash With FluentD



# Logstash vs FluentD



Features	Logstash	FluentD
Language	Written in JRuby; requires a Java runtime	Written in C and Ruby; no Java dependencies
Ecosystem and Plugins	200 plugins; centralized by Elastic	500 plugins; decentralized in various repositories
Data Transport	No built-in buffering; uses external queues like Redis or Kafka for persistence	Built-in buffering system; reliable without external dependencies
Performance	Consumes more memory than FluentD	More memory-efficient than Logstash
Event Routing	Generally higher due to durability features and multi-AZ deployments	Uses tagging for routing; simpler and more straightforward
Log Parsing	Uses if-then-else statements for routing; more complex configuration	Built-in parsers for formats (JSON, regex, CSV, etc.); less need for external plugins

© Copyright KodeKloud

## Logstash

**Language:** Written in JRuby, requiring a Java runtime on the host machine.

**Ecosystem and Plugins:** Has around 200 plugins available in a centralized repository managed by Elastic.

**Data Transport:** Lacks an built-in buffering system, relying on external queues like Redis or Kafka for data persistence across restarts.

**Performance:** Generally consumes more memory than FluentD.

**Event Routing:** Uses if-then-else conditional statements for routing events, which can be more complex to configure.

**Log Parsing:** Relies on plugins for log parsing, requiring additional configuration.

plugins.

# FluentD vs FluentBit



Criteria	FluentD	FluentBit
Resource Consumption	Heavier resource usage; requires more memory and CPU; suitable for server environments	Lightweight; minimal resource usage; ideal for edge and embedded systems
Deployment and Usage	More complex setup; requires Ruby runtime; rich feature set and plugins	Easier to deploy with a smaller binary size; quick setup with minimal configuration
Performance	Generally slower due to additional features and Ruby-based implementation	High performance with low-latency data collection and forwarding
Configuration	More complex with multiple files; supports advanced features like routing, buffering, and retry logic	Simpler; uses a single configuration file with straightforward syntax
Data Processing Capabilities	Advanced features for parsing, flexible routing, and data transformation	Basic capabilities for filtering, parsing, and buffering
Use Cases	Ideal for centralized logging, complex log aggregation, and extensive customization	Ideal for edge computing, IoT devices, and scenarios needing lightweight, high-performance log forwarding

© Copyright KodeKloud

## FluentD vs Fluent Bit

### 1. Origin and Development

1. **FluentD:** Developed by Treasure Data, now a part of the Cloud Native Computing Foundation (CNCF).
2. **Fluent Bit:** Also developed by the Fluent community, designed as a lightweight and high-performance log processor and forwarder.

### 2. Performance

1. **FluentD:** More resource-intensive due to its feature set and design, which can be demanding on memory and

CPU.

2. **Fluent Bit**: Optimized for performance and low resource usage, making it suitable for edge devices and high-volume log processing.

## 1. Language and Architecture

1. **FluentD**: Written in C and Ruby, leveraging Ruby for its plugin ecosystem and flexibility.
2. **Fluent Bit**: Written in C, designed to be lightweight and efficient, with a minimalistic architecture aimed at high throughput.

## 2. Plugins and Extensibility

1. **FluentD**: Offers a rich ecosystem of plugins (over 900), covering a wide range of inputs, filters, and outputs, allowing extensive customization.
2. **Fluent Bit**: Has a smaller but growing set of plugins (over 150), focusing on essential functionalities and integration with cloud-native environments.

## 3. Configuration

1. **FluentD**: Uses a configuration file format based on a JSON-like syntax, which some users find complex, especially for large configurations.
2. **Fluent Bit**: Employs a straightforward configuration syntax, often simpler and easier to manage, making it user-friendly for basic and advanced setups.

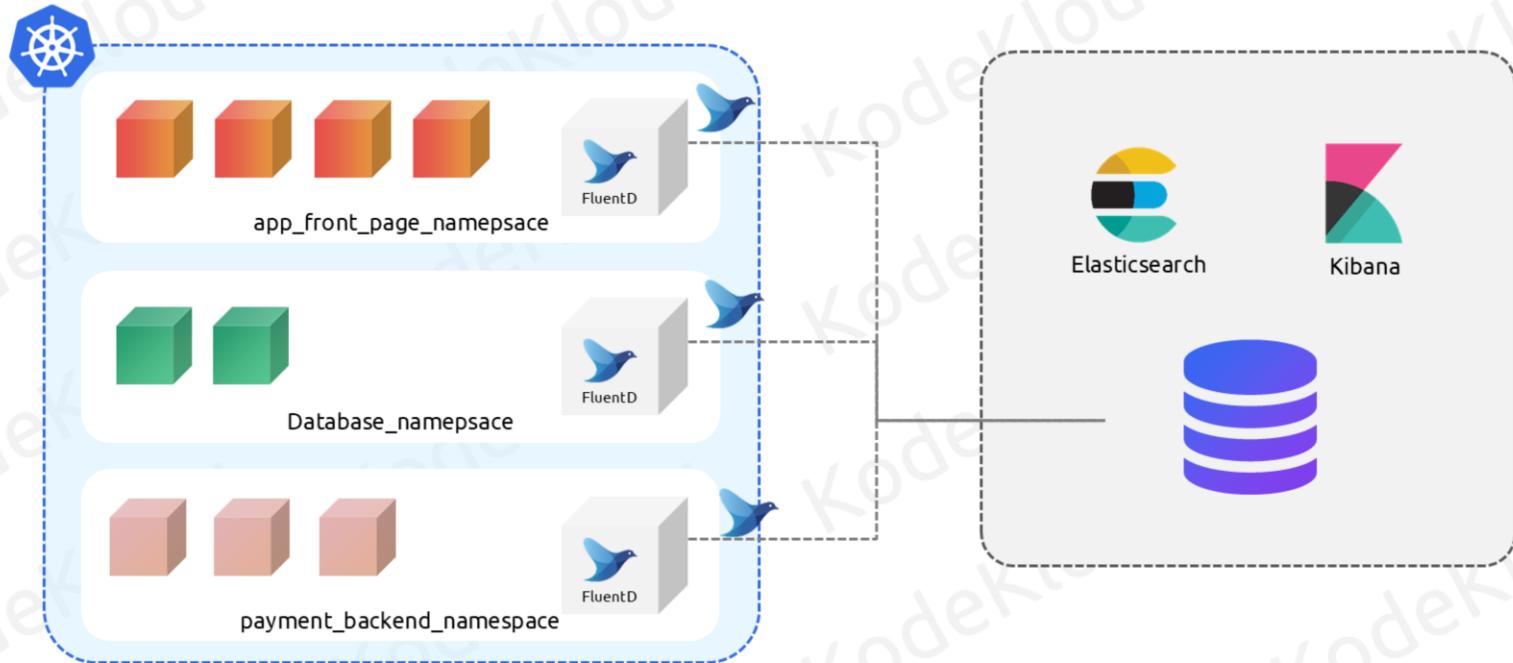
## 4. Deployment and Use Cases

1. **FluentD**: Typically used in environments where complex log processing and transformation are required, especially within the Elastic ecosystem.
2. **Fluent Bit**: Designed for lightweight log forwarding and collection, ideal for edge computing, IoT, and as a log processor for Kubernetes and cloud-native setups.

## 5. Community and Ecosystem

1. **FluentD**: Strong community support with extensive documentation and integration within the Kubernetes ecosystem and other cloud services.
2. **Fluent Bit**: Growing community with a focus on performance and simplicity, gaining traction in the cloud-native space for its efficiency and ease of use.

# FluentBit's Role as a Lightweight Log Shipper



# Input, Filter, and Output Plugins in FluentBit



# Input Plugins

**Purpose:** Input plugins collect data from various sources.

## tail

Reads data from a text file as it is written (tailing a log file).

## systemd

Collects logs from the systemd journal.

## tcp

Listens for TCP connections and reads the data sent to it.

```
[INPUT]
Name tail
Path /var/log/nginx/access.log
Tag nginx.access
Parser nginx
```

```
[INPUT]
Name systemd
Tag host.*
```

```
[INPUT]
Name tcp
Listen 0.0.0.0
Port 5170
Tag tcp.input
```

# Filter Plugins

**Purpose:** Filter plugins modify, enrich, or drop records based on specific rules.

## grep

Filters records by matching a regular expression.

```
[FILTER]
Name grep
Match nginx.access
Regex message error
```

## modify

Adds, removes, or modifies fields in the log records.

```
[FILTER]
Name modify
Match *
Add service nginx
```

## parser

Parses specific fields in the logs.

```
[FILTER]
Name parser
Match nginx.access
Key_Name message
Parser json
```

# Output Plugins

**Purpose:** Output plugins send the collected and processed data to various destinations.



Sends data to Elasticsearch.

```
[OUTPUT]
Name es
Match *
Host 127.0.0.1
Port 9200
Index fluentbit
Type _doc
```



Sends data to an HTTP endpoint.

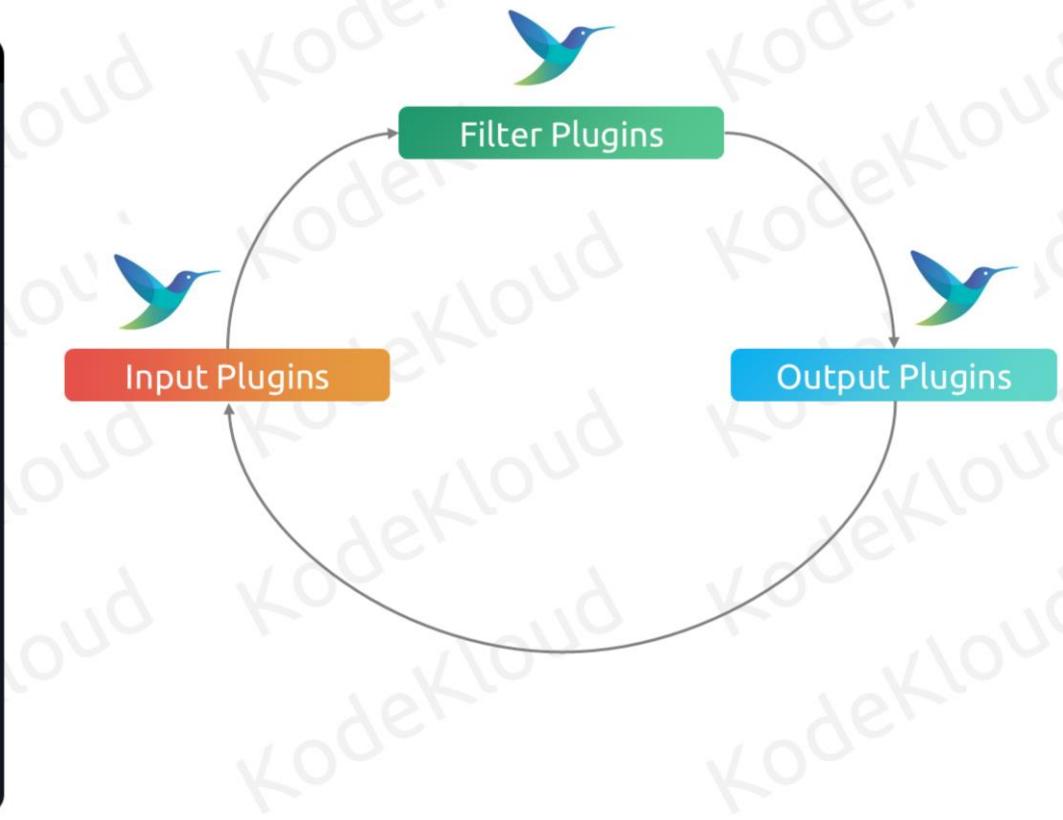
```
[OUTPUT]
Name http
Match *
Host example.com
Port 80
URI /data
Format json
```

# Putting It All Together

```
Input section
[INPUT]
 Name tail
 Path /var/log/nginx/access.log
 Tag nginx.access
 Parser nginx

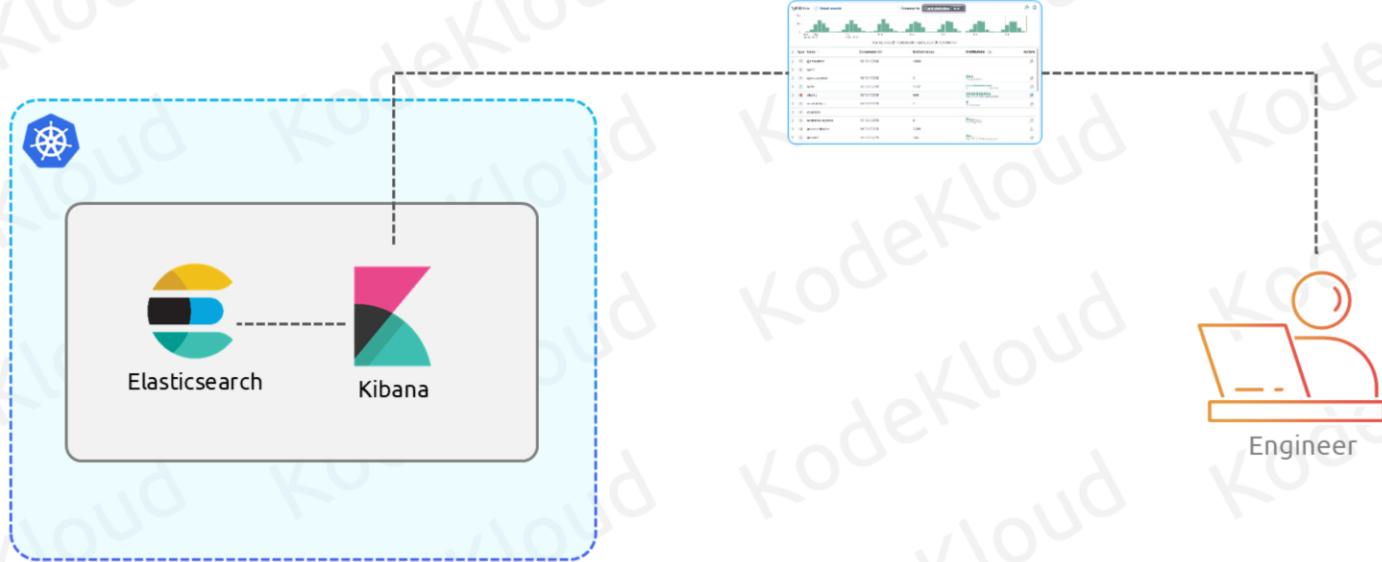
Filter section
[FILTER]
 Name grep
 Match nginx.access
 Regex message error

Output section
[OUTPUT]
 Name es
 Match *
 Host 127.0.0.1
 Port 9200
 Index fluentbit
 Type _doc
```



# **Elasticsearch and Kibana Deployment on Kubernetes**

# Elasticsearch and Kibana on Kubernetes



**Demo**

Deploy Elasticsearch  
on Kubernetes

# Demo

## Deploy Kibana on Kubernetes

# Demo

Resource allocation (CPU, memory) for Elasticsearch and Kibana

# Security Considerations for Elasticsearch and Kibana on Kubernetes

1 | Secure Kubernetes Cluster

2 | Enable SSL/TLS Encryption

3 | Elastic Stack Security Features

4 | Restrict Elasticsearch Access

```
networkPolicy:
 - name: allow-kibana
 podSelector:
 matchLabels:
 app: kibana
 ingress:
 - from:
 - podSelector:
 matchLabels:
 app: kibana
 egress:
 - to:
 - podSelector:
 matchLabels:
 app: elasticsearch
```

© Copyright KodeKloud

## Secure Kubernetes Cluster

Remember that the Elasticsearch cluster is only as secure as the underlying Kubernetes cluster, so ensure the Kubernetes environment is properly secured to prevent unauthorized access.

## Enable SSL/TLS Encryption

Use SSL/TLS encryption to secure traffic between browsers and Kibana server, preventing eavesdropping and tampering.

## Elastic Stack Security Features

ensuring that only necessary users can access sensitive data.

# Security Considerations for Elasticsearch and Kibana on Kubernetes

1 | Secure Kubernetes Cluster

2 | Enable SSL/TLS Encryption

3 | Elastic Stack Security Features

4 | Restrict Elasticsearch Access

```
Enable SSL/TLS encryption for
Kibana
kibana:
 server:
 ssl:
 enabled: true
 certificate: /path/to/cert.pem
 key: /path/to/key.pem
```

# Security Considerations for Elasticsearch and Kibana on Kubernetes

1 | Secure Kubernetes Cluster

2 | Enable SSL/TLS Encryption

3 | Elastic Stack Security Features

4 | Restrict Elasticsearch Access

```
Enable Elastic Stack security
features
elasticsearch:
 security:
 enabled: true
 authc:
 realms:
 - type: basic
 order: 0
 authz:
 roles:
 - type: basic
 order: 0
```

# Security Considerations for Elasticsearch and Kibana on Kubernetes

1 | Secure Kubernetes Cluster

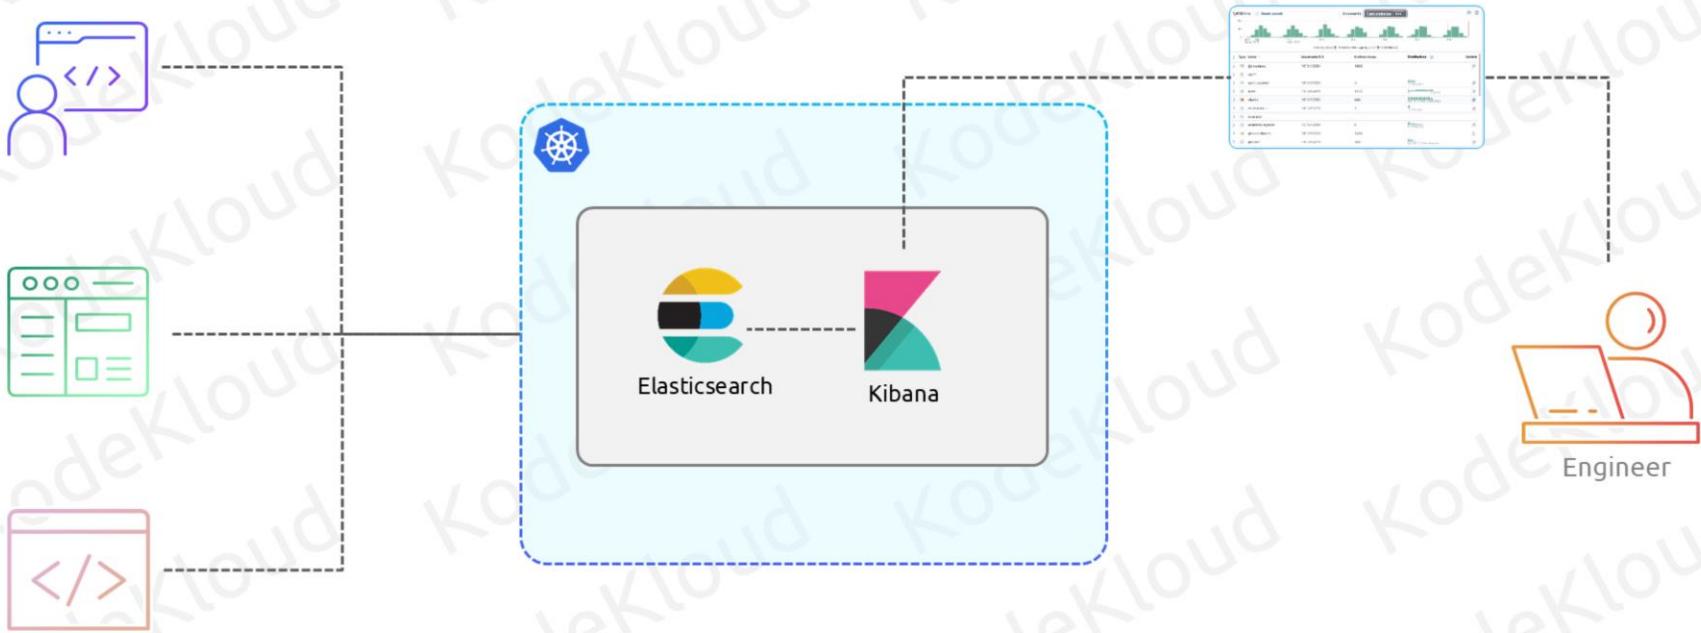
2 | Enable SSL/TLS Encryption

3 | Elastic Stack Security Features

4 | Restrict Elasticsearch Access

```
Restrict access to Elasticsearch
by only allowing traffic from
Kibana
elasticsearch:
 network:
 host: kibana
 port: 9200
```

# Scaling Elasticsearch Nodes



# Scaling Elasticsearch Nodes

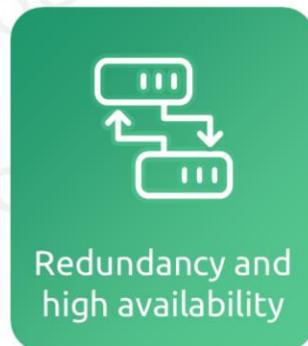
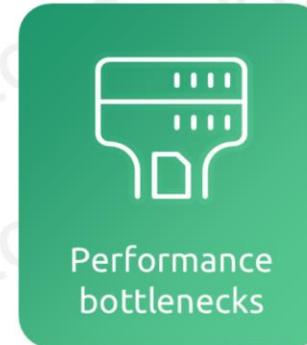


When should we consider scaling Elasticsearch nodes?

© Copyright KodeKloud

When Should We Consider Scaling Elasticsearch Nodes?

# Scaling Elasticsearch Nodes



© Copyright KodeKloud

## Increased Data Volume

Explanation: When the amount of stored data grows, more nodes are needed to handle the increased storage requirements efficiently.

## High Query Load

Explanation: To maintain performance during high search or query load, adding nodes can distribute the query processing

workload.

#### Performance Bottlenecks

Explanation: When nodes are frequently experiencing high CPU or memory usage, scaling can alleviate these bottlenecks.

#### High Ingestion Rates

Explanation: With an increased rate of data ingestion, additional nodes help manage and index the incoming data more effectively.

#### Redundancy and High Availability

Explanation: To ensure data redundancy and high availability, scaling nodes provides more replicas and fault tolerance.

#### Geographical Distribution

Explanation: For globally distributed applications, scaling nodes across different regions can reduce latency and improve access speed.

# Scaling Elasticsearch Nodes



© Copyright KodeKloud

What are some ways recommended to scale Elasticsearch nodes?

What Are Some Ways Recommended to Scale Elasticsearch Nodes?

# Scaling Elasticsearch Nodes



© Copyright KodeKloud

01

Vertical Scaling

02

Horizontal Scaling

03

Node Types Specialization

04

Sharding

05

Index Lifecycle Management (ILM)

06

Snapshot and Restore

## Vertical Scaling

Explanation: Increasing the resources (CPU, memory, storage) of existing nodes to handle more data and queries.  
Horizontal Scaling

Explanation: Adding more nodes to the cluster to distribute the workload and increase capacity.  
Node Types Specialization

Explanation: Using specialized nodes for different roles such as master, data, and ingest nodes to optimize cluster performance.

### Sharding

Explanation: Dividing indices into smaller pieces (shards) that can be distributed across multiple nodes to balance the load.

### Index Lifecycle Management (ILM)

Explanation: Implementing ILM to automate the management of indices through different phases, reducing the load on nodes.

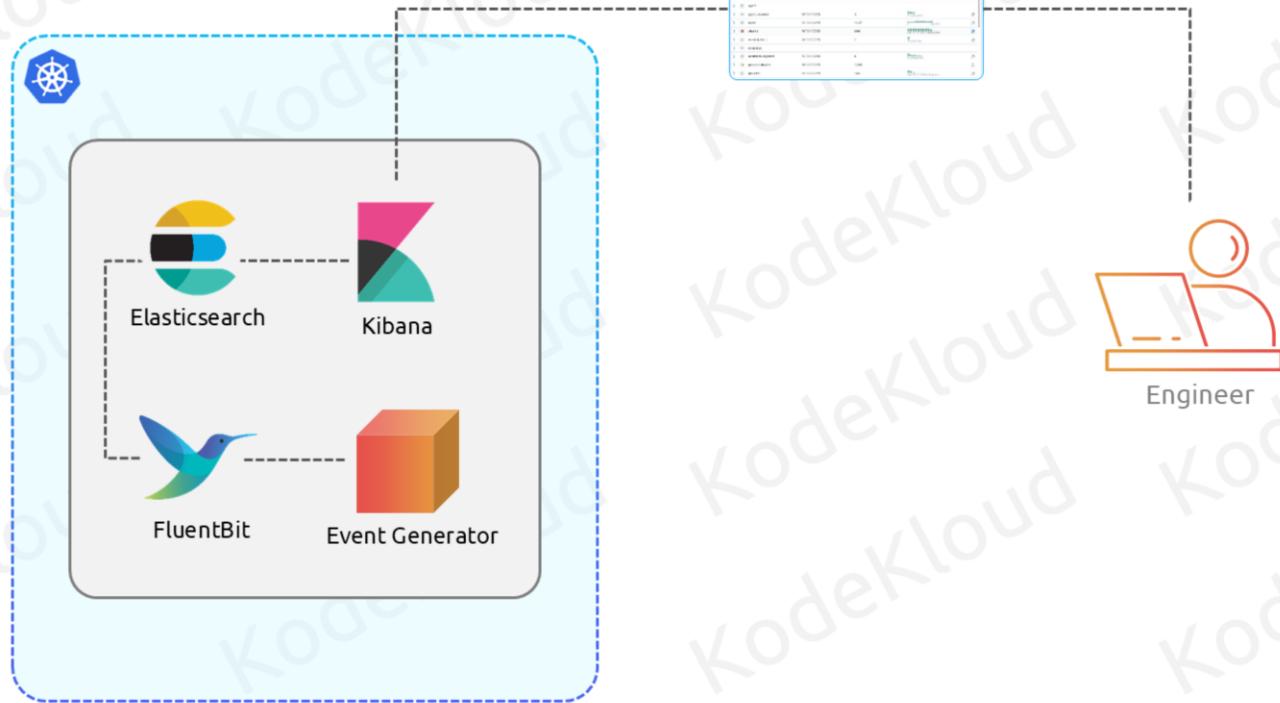
### Snapshot and Restore

Explanation: Using snapshots for backup and disaster recovery, enabling easy restoration and scaling of nodes when needed.

These sections and sub-points can be elaborated into detailed explanations or visualizations on your slide to help your audience understand when and how to scale Elasticsearch nodes effectively.

# **Deploying Mock Ecommerce Application on K8**

# Deploying Mock Ecommerce Application and Kibana on K8s



# Demo

## Setting up Event Generator App

# Demo

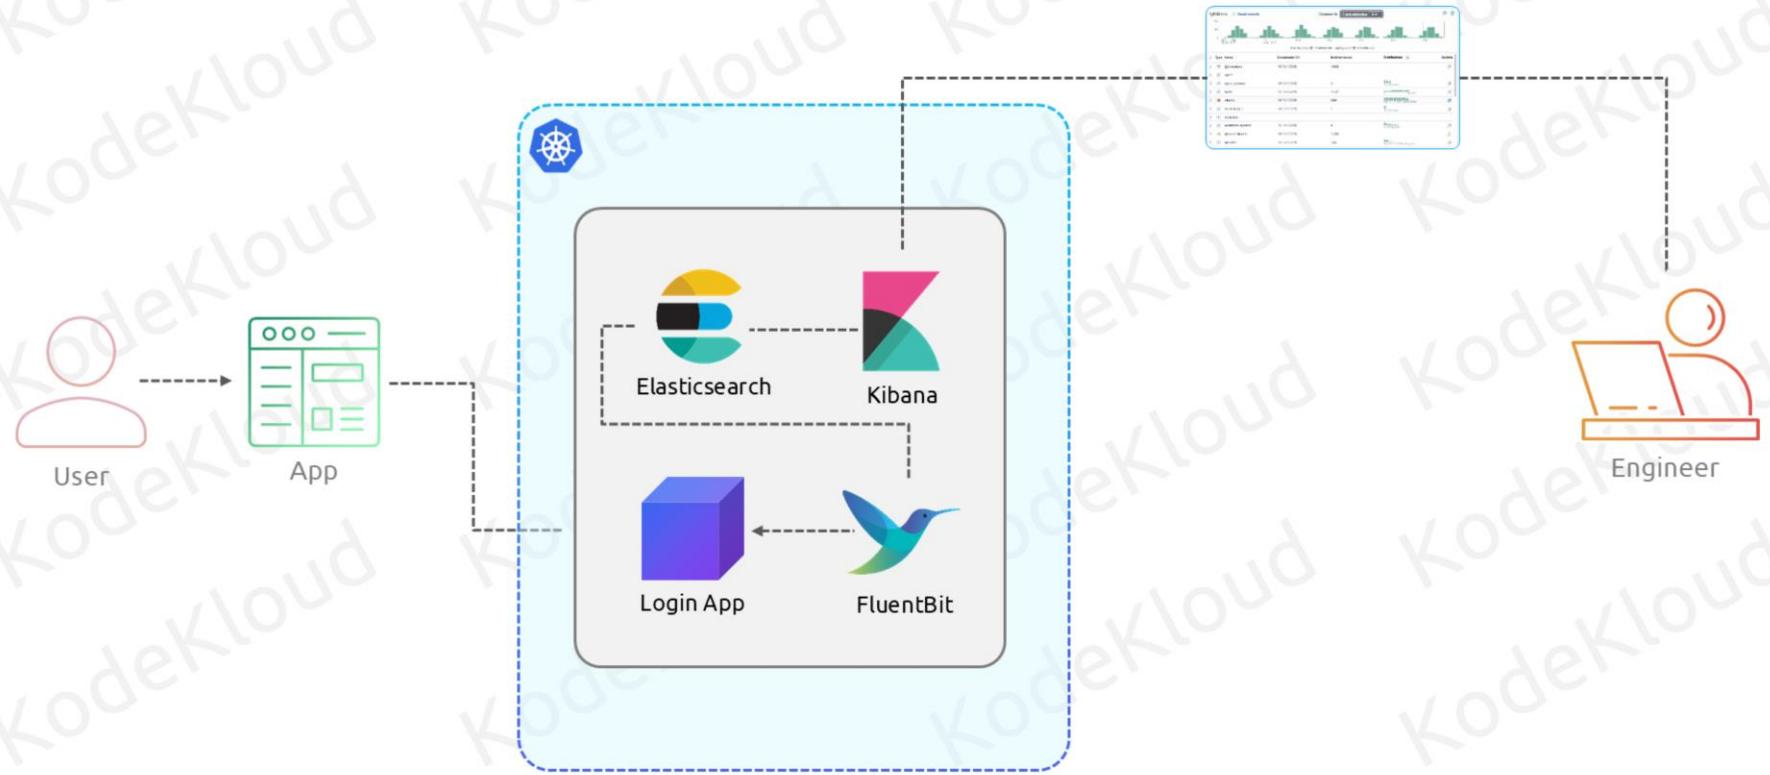
Setting up FluentBit to  
monitor Application Logs

**Demo**

Exploring Application  
Logs in Kibana

# **Monitoring a Python Application With FluentBit and Elasticsearch**

# Deploying Login App and Kibana on Kubernetes



© Copyright KodeKloud

# Demo

Deploying and  
validating the login  
app on Kubernetes

# Demo

Configuring FluentBit  
to collect Python  
Application Logs

# Demo

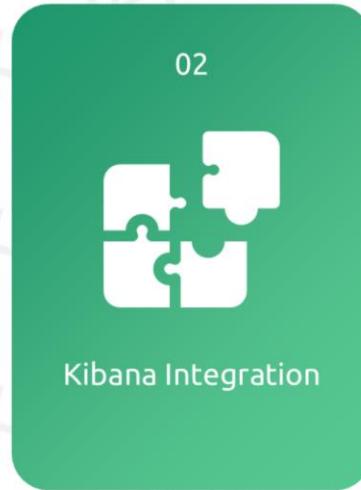
Building Kibana  
Dashboards to visualize  
our application

# Elastic Cloud

# Elastic Cloud – Offerings



01  
Managed Elasticsearch Service



02  
Kibana Integration



03  
Security and Compliance



04  
Multi-Cloud and Hybrid Deployment

© Copyright KodeKloud

## 1. Managed Elasticsearch Service

Elastic Cloud offers a fully managed Elasticsearch service, allowing users to deploy, manage, and scale Elasticsearch clusters without worrying about the underlying infrastructure. This service includes automatic updates, backups, and monitoring.

## 2. Kibana Integration

Elastic Cloud integrates seamlessly with Kibana, providing powerful visualization and analysis tools for Elasticsearch data. Users can create interactive dashboards, perform real-time analysis, and gain insights from their data effortlessly.

## 3. Security and Compliance

Elastic Cloud provides robust security features, including encryption at rest and in transit, role-based access control, and single sign-on (SSO) capabilities. It also complies with various industry standards and regulations, ensuring data protection and privacy.

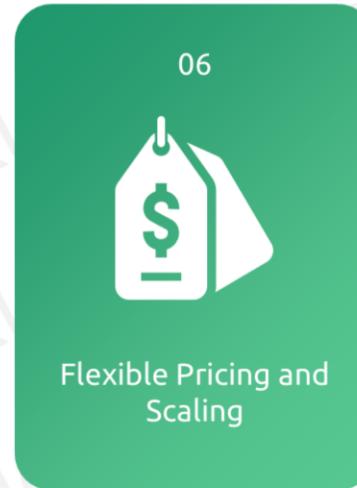
#### **4. Multi-Cloud and Hybrid Deployment**

Elastic Cloud supports multi-cloud and hybrid deployment options, allowing users to deploy their Elasticsearch clusters on their preferred cloud providers, such as AWS, Azure, and Google Cloud, or on-premises.

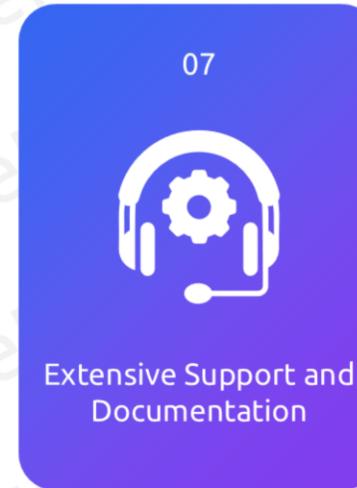
# Elastic Cloud – Offerings



Advanced Features and  
Machine Learning



Flexible Pricing and  
Scaling



Extensive Support and  
Documentation

© Copyright KodeKloud

## 5. Advanced Features and Machine Learning

Elastic Cloud offers advanced features like machine learning, anomaly detection, and alerting. These features help users to detect unusual patterns, forecast trends, and receive notifications about critical events in their data.

## 6. Flexible Pricing and Scaling

Elastic Cloud provides flexible pricing plans based on the resources used, making it cost-effective for various use cases. Users can scale their clusters up or down based on their needs, ensuring they only pay for what they use.

## 7. Extensive Support and Documentation

Elastic Cloud offers extensive support options, including 24/7 technical support, community forums, and comprehensive documentation. Users have access to resources and assistance to ensure their deployments run smoothly and efficiently.

<https://cloud.elastic.co/pricing>

**Demo**

Setting up free  
Elastic Cloud account

# Demo

Monitoring  
Kubernetes Cluster  
using ElasticAgent



# KodeKloud

© Copyright KodeKloud

Follow us on <https://kodekloud.com/> to learn more about us.